

# Distributed Systems

## Lecture 1

### Introduction and Review

Joseph Phillips

Copyright (c) 2019

Last modified 2019 September 12

# Topics

- Motivations
- Issues
- Low level I/O in Unix/C
- Low level I/O in Java
- Support for the class

# Motivation

- Why distribute a system?
- Disadvantages:
  - Reliance on a network
  - Lack of a common clock
  - Inherently less stable
  - Inherently less secure



# Advantage #1

- Ability to use non-local service
  - ssh into remote computer
  - ftp/sftp into remote computer
  - http into remote web-server

# Advantage #2



- Ability to scale solutions
  - Distributed computation
    - Notable example: [SETI@home](#)
    - Many others:  
[https://en.wikipedia.org/wiki/List\\_of\\_distributed\\_computing\\_projects](https://en.wikipedia.org/wiki/List_of_distributed_computing_projects)
  - Distributed responsibilities
    - Login server
    - Database server

# Advantage #3

- A different form of robustness
  - Use network to our advantage
  - If machine goes down, others still work

# Advantage #4

- Pay for computing services only when they are needed
  - “Replicated load-balancing service” design pattern from Brendan Burns



# Advantage #5

- Implementation language independence
  - Similar notion as “interfaces” in object-oriented programming
- You use your language, I’ll use mine
  - Your instructor is fond of C/C++, you may prefer Java or Python
- Some languages are better for some tasks
  - Fortran for matrix operations
  - Lisp for some A.I. applications
  - C for systems
  - Java or Python for getting something robust working quickly (programmer time)
- Support legacy systems / legacy languages
  - Cobol is still used

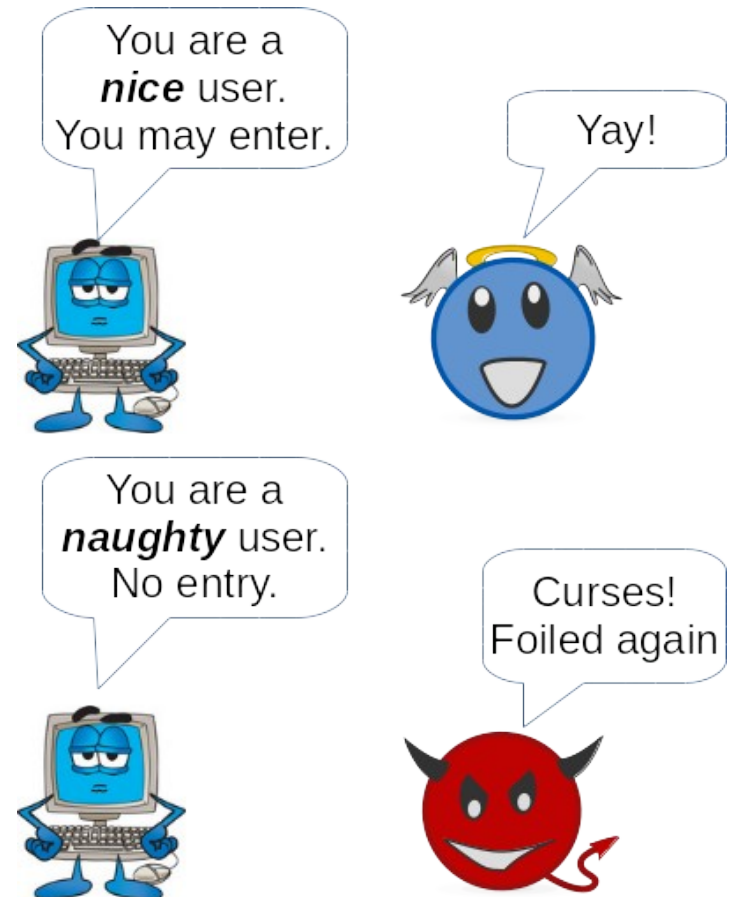


# Issues: Protocols

- Internet Layer Protocols
  - IP: Internet Protocol
  - vs the earlier Network Control Protocol, a point-to-point protocol
- Transport Layer Protocols:
  - TCP: Transmission Control Protocol
  - Vs UDP (User Datagram Protocol)
- Application Layer Protocols:
  - SSH: Secure SHell
  - SNMP: Simple Network Management Protocol
- Above the application layer
  - XML
  - Json

# Issues: control and access


- May want to only allow
  - Certain users
  - From certain machines



# Review: I/O in Unix/C

- File descriptors
  - Integers that refer to an index in an array of open files

Integer value	Symbolic constant	FILE* equivalent
0	STDIN_FILENO	stdin
1	STDOUT_FILENO	stdout
2	STDERR_FILENO	stderr



# File descriptors for files

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```

- For read-only

```
int open(const char *pathname, int flags, mode_t mode);
```

- For read, write, read/write

```
int creat(const char *pathname, mode_t mode);
```

- For write-only

- Flags:

- One of `O_RDONLY`, `O_WRONLY`, `O_RDWR`

- Bitwise OR-ed with one or more of `O_CREAT`, `O_TRUNC`, `O_APPEND` (and some more)

- mode is permissions

- If return -1 then error

# write()

```
int write(int fd, char* bufferPtr,  
          size_t numBytes)
```

- Writes numBytes pointed to by bufferPtr to fd.
- Returns number of bytes written, or -1 on error.

# read()

- `int read(int fd, char* bufferPtr, size_t bufferSize)`
  - Reads up to `bufferSize` bytes from `fd` and puts them into `bufferPtr`.
  - Returns number of bytes read from file, either
    - 0 (*"No more left!"*),
    - `bufferSize` (*"Here's a whole buffer full!"*),
    - somewhere inbetween (*"Here's all that's left"*), or,
    - -1 (*"Error!"*)
- But, for network I/O, please use . . .

# rio\_read()

```
// PURPOSE: To be the "robust" version of
read()
//   in the manner advocated by Bryant and
//   O'Hallaron. 'fd' tells the file descriptor from
//   which to read. 'usrbuf' tells the buffer into
//   which to read. 'n' tells the length of 'usrbuf'.
//   Returns number of bytes read or -1 on
//   some type of error.
```

```
ssize_t rio_read (int  fd,
                  char* usrbuf,
                  size_t n
                  )
```

```
{
    // I. Application validity check:
```

```
    // II. Attempt to read:
```

```
    ssize_t  nread;
    size_t    nleft = n;
    char*     bufp  = usrbuf;
```

```
    while (nleft > 0)
    {
        if ((nread = read(fd, bufp, nleft)) < 0)
        {
            if (errno == EINTR) // interrupt by sig handler?
                nread = 0;      // yes: read() again
            else
                return -1;      // no: errno set by read()
        }
        else
            if (nread == 0)
                break;          // EOF

        nleft -= nread;
        bufp += nread;
    }

    // III. Finished:
    return (n - nleft);        // return >= 0
}
```

# Your Turn!


Write a program to copy a file  
given paths on the command line:

```
$ ourCopy sourceFile destFile
```



# Remember:

- In Unix/C: `fork()` copies the file descriptor table
- A pipe or socket may still be open if only one process `close()`s it
- Therefore, immediately `close()` the file descriptor you won't use



# *Quick!*

## What's wrong with this program?

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>

const int BUFFER_LEN = 256;

#define TEXT "Something to send down the pipe."

int main ()
{
    int fd[2];

    if (pipe(fd) < 0)
        exit(EXIT_FAILURE);

    pid_t childId = fork();

    if (childId < 0)
        exit(EXIT_FAILURE);

    if (childId == 0)
    {
        char buffer[BUFFER_LEN];
        int numBytes;

        if (read (fd[0],buffer,BUFFER_LEN) > 0)
        {
            printf("Child received \"%s\"\n",buffer);
        }

        close(fd[0]);
    }
    else
    {
        printf("Parent sending \"%s\"\n",TEXT);
        write(fd[1],TEXT,sizeof(TEXT));
        close(fd[1]);
        wait(NULL);
    }

    return(EXIT_SUCCESS);
}
```

# Java streams

- Input and Output
- Often filters are chained together
- Filters for
  - buffering (efficiency)
  - compression/uncompression
  - encrypting/decrypting

# Java output streams

- From java.io.OutputStream
  - public abstract class OutputStream
  - public abstract void write (int b) throws IOException
    - Writes lowest byte of integer 'b'
  - public void write (byte[ ] data) throws IOException
  - public void write (byte[ ] data, int offset, int length) throws IOException
    - Write array (or subarray) of bytes
  - public void flush () throws IOException
  - public void close () throws IOException

# OutputStream Example

```
public static void generateChars    (OutputStream out)
    throws IOException
{
    int firstPrintChar = 33;
    int numPrintChars  = 94;
    int numCharsPerLine = 72;
    int start          = firstPrintChar;

    while (true)
    {
        for (int i = start; i < start + numCharsPerLine; i++)
            out.write((i-firstPrintChar) % numPrintChars + firstPrintChar);
        out.write('\r');
        out.write('\n');
        start = ((start+1) - firstPrintChar) % numPrintChars + firstPrintChar;
    }
}
```

# But wait! Is that efficient?

- Sending bytes one-at-a-time is inefficient
  - Smallest TCP header 40 bytes.
  - Sending 1 application byte really means sending at least 40 more
  - Thus sending 100 bytes one-at-a-time could mean sending 4100 bytes total
- Better to buffer bytes into an array . . .

# Improved OutputStream Example

```
// From Eliot Rusty Harold
// "Java Network Programming"
// O'Reilly, 2014
import java.io.OutputStream;
import java.io.IOException;

public class GenChars
{
    public static void generateChars    (OutputStream out)
        throws IOException
    {
        int firstPrintChar = 33;
        int numPrintChars  = 94;
        int numCharsPerLine = 72;
        int start          = firstPrintChar;

        byte[] line        = new byte[numCharsPerLine + 2];

        while (true)
        {
            for (int i = start; i < start + numCharsPerLine; i++)
                line[i-start] = (byte)
                    ((i-firstPrintChar) % numPrintChars +
firstPrintChar);
```

```
                line[numCharsPerLine+0] = (byte)'\r';
                line[numCharsPerLine+1] = (byte)'\n';
                out.write(line);
                start = ((start+1) - firstPrintChar) %
numPrintChars + firstPrintChar;
            }
        }

    public static void main(String[] args)
    {
        try
        {
            generateChars(System.out);
        }
        catch (IOException except)
        {
        }
    }
}
```

# Is buffering the end of our troubles?

- Not quite!
- `flush()` the output whenever you want to send it immediately
- Be sure to `close()` streams at end, even upon exceptions
  - Otherwise could memory leak file descriptors

```
OutputStream out = null;
try
{
    out = new FileOutputStream("out.txt");
    // work with out
}
catch (IOException ex)
{
    System.err.println(ex.getMessage());
}
finally
{
    if (out != null)
    {
        try
        {
            out.close();
        }
        catch (IOException ex)
        {
            // Ignore
        }
    }
}
```



# InputStream

- public abstract class InputStream
- public abstract int read () throws IOException
  - Reads single byte
  - -1 means “no more”
  - Blocks (waits) until byte is available
- public int read (byte[ ] input) throws IOException
- public int read (byte[ ] input, int offset, int length) throws IOException
  -
- public long skip (long n) throws IOException
- public int available () throws IOException
  - Tells how many bytes can be read without blocking
- public void close() throws IOException

# To naively read one byte a time

```
byte [ ] input = new byte[10];  
for (int i = 0; i < input.length; i++)  
{  
    int b = in.read();  
    if (b < 0) break;  
    input[i] = (byte)b;  
}
```

# More efficient

```
int bytesRead = 0;
int bytesToRead = 64;
byte [ ] input = new byte[bytesToRead];
while (bytesRead < bytesToRead)
{
    int result = in.read(input,
bytesRead,bytesToRead - bytesRead);
    if (result < 0) break;
    bytesRead += result;
}
```

# File Streams

- **FileInputStream:**
  - `FileInputStream(File file)`
  - `FileInputStream(FileDescriptor fdObj)`
  - `FileInputStream(String name)`
- **FileOutputStream:**
  - `FileOutputStream(File file)`
  - `FileOutputStream(File file, boolean append)`
  - `FileOutputStream(FileDescriptor fdObj)`
  - `FileOutputStream(String name)`
  - `FileOutputStream(String name, boolean append)`

# Your Turn!

Write a program to copy a file  
given paths on the command line:

```
$ java CopyFile CopyFile.java CopyFileCopy.java
```

# Your Turn! (continued)

```
import java.io.InputStream;
import java.io.FileInputStream;
import java.io.OutputStream;
import java.io.FileOutputStream;
import java.io.IOException;

public class CopyFile
{
    public static void main (String[] args)
    {
        if (args.length < 2)
        {
            System.err.println("Usage:\tjava CopyFile
<sourceFile> <destFile>");
            return;
        }

        copy(args[0],args[1]);
    }
}
```

```
public static void copy
    (String sourcePath,
     String destPath
    )
{
    final int MAX_BUFFER_LEN = 256;
    FileOutputStream out = null;
    FileInputStream in = null;
    byte[] byteArray = new
byte[MAX_BUFFER_LEN];

    // YOUR CODE HERE
}
}
```

# FilterStreams and Chaining

- FilterStreams come in two types:
  - input
  - output
- They do processing on their streams
  - output stream compresses/input stream uncompresses
  - output stream encrypts/input stream decrypts
- They are meant to be chained together, e.g.  

```
InputStream in = new FileInputStream("data.txt");  
in = new BufferedInputStream(in);  
  
// or  
  
InputStream in = new BufferedInputStream(new FileInputStream("data.txt"));
```
- Buffered streams are a special case:
  - `public BufferedInputStream(InputStream in)`
  - `public BufferedInputStream(InputStream in, int bufferSize);`
  - `public BufferedOutputStream(OutputStream in)`
  - `public BufferedOutputStream(OutputStream in, int bufferSize);`

# PrintStream

- Please do NOT use PrintStream over the network!
  - How say “next line”?
    - Unix: '\n'
    - Windows: '\r' '\n'
    - MAC: '\r'
    - println() on local system will be fine, but maybe not on remote system
  - Similarly, char encoding on local system may be different than remote (e.g. UTF-8, UTF-16, CP1252 (US Windows), SJIS (Japan), etc.)
  - Ignores exceptions, but these might be important in network communication



# DataOutputStream

public final void writeBoolean (boolean b) throws  
IOException

public final void writeByte (int b) throws IOException

public final void writeShort (int s) throws IOException

public final void writeChar (int c) throws IOException

public final void writeInt (int i) throws IOException

public final void writeLong (long l) throws IOException

public final void writeFloat (float f) throws IOException

public final void writeDouble (double d) throws IOException

public final void writeChars (String s) throws IOException

public final void writeBytes (String s) throws IOException

public final void writeUTF (String s) throws IOException

# DataInputStream

```
public final boolean readBoolean () IOException
public final byte readByte () IOException
public final char readChar () IOException
public final short readShort () IOException
public final int readInt () IOException
public final long readLong () IOException
public final float readFloat () IOException
public final double readDouble () IOException
public final String readUTF () IOException
```

# DataStream

- Meant to transfer data to another system that understands Java datatypes
  - Another Java program
  - A program (perhaps C) that uses Java types
- Java types
  - Big (standard network) endian
  - 2's complement integers
    - boolean, byte = 1 byte
    - short = 2 bytes
    - int, float = 4 bytes
    - long, double = 8 bytes
    - char = 2 unsigned bytes
  - writeChars(): UTF-16
  - writeBytes(): low byte only
  - writeUTF(): Java's ***variant*** of UTF-8 (only use with other Java programs)

# Readers and Writers

- Streams are for reading/writing ***buffers*** and ***data***
  - Analogous to ***file descriptors*** in C
- Readers and Writers specialize in ***text***
  - Analogous to ***FILE\**** in C
- By default reads/writes unsigned 16-bit chars (UTF-16)

# Writers

- protected Writer()
- protected Writer (Object lock)
- public abstract void (char[ ] text, int offset, int length) throws IOException
- public void write (int c) throws IOException
- public void write (char[ ] text) throws IOException
- public void write (String s) throws IOException
- public void write (String s, int offset, int length) throws IOException
- public abstract void flush () throws IOException
- public abstract void close () throws IOException

# Readers

- protected Reader ()
- protected Reader (Object lock)
- public abstract int read (char[ ] text, int offset, int length)
- public int read () throws IOException
- public int read (char[ ] text) throws IOException
- public long skip (long n) throws IOException
- public boolean ready ()
- public boolean markSupport()
- public void mark (int readAheadLimit) throws IOException
- public void reset () throws IOException
- public abstract void close() throws IOException

# Getting a Unix System

- Do you already have your own Linux, Mac OS or Unix machine? Awesome! Use it.
- See Prof Joe after class to get an account on the virtual machine:  
`phillips373.cdm.depaul.edu`.

# References

- M. Tim Jones “*GNU/Linux Application Programming, 2nd Ed*” Course Technology, Cengage Learning. 2008
- Eliot Rusty Harold “*Java Network Programming, 4th Ed*” O'Reilly. 2014
- Brendan Burns “*Designing Distributed Systems: Patterns and Paradigms for Scalable, Reliable Services*” O'Reilly. 2018