

Distributed Systems

Lecture 2

Distributed Systems *Lite*:
Multiprocessing on a single machine


Joseph Phillips
Copyright (c) 2019
Last modified 2019 September 24

Topics

- Motivation
- Issues
- Inter-process communication
 - pipes
 - socketpairs
 - messages
 - shared memory

Motivation

Question: We have the Internet! Why limit ourselves to one machine?




Motivation

- Answer #1: To debug a system before distributing it
- Answer #2: To make sure even a distributed system will work even if there is no network

Motivation

Question: Fine! Limit yourself to one machine. If you want to do two things at once, just use threads instead of inter-process communication



Motivation

- `fork()/execl()` is more flexible than threads
 - Child process can run another program without destroying parent
- `fork()/execl()` is more robust than threads
 - If child crashes, does not crash parent

Issues

- Do not have to worry about networking (yet!)
- Do have to worry about
 - Inter-process communication

Inter-process communication: pipes (1)

- Pipes:

- One way, file-descriptor based communication between related processes
 - child-to-parent
 - parent-to-child
 - child-to-child
- Takes array of 2 ints
 - array[0]: reading fd (think STDIN_FILENO = 0)
 - array[1]: writing fd (think STDOUT_FILENO = 1)

- pipe(int[2])

- pipe2(int[2],int flags)

- flags are bitwise OR of:
 - O_NONBLOCK: make reading non-blocking
 - O_CLOEXEC: for use in some multi-threaded apps

- **Remember! Processes should close unneeded end!**

```
#include <unistd.h>
...
const int PIPE_READ = 0;
const int PIPE_WRITE= 1;
int      myPipe[2];

if (pipe(myPipe) == 0)
{
    char myArray[6];
    write(myPipe[PIPE_WRITE], "Hello!", 6);
    read (myPipe[PIPE_READ ], myArray, 6);
}
```



myPipe: An OS-owned buffer

Inter-process communication: pipes (2)

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>

const int BUFFER_LEN = 256;

#define TEXT \
    "Something to send down the pipe."

int main ()
{
    int fd[2];

    if (pipe(fd) < 0)
        exit(EXIT_FAILURE);

    pid_t childId = fork();

    if (childId < 0)
        exit(EXIT_FAILURE);
```

```
    if (childId == 0)
    {
        char buffer[BUFFER_LEN];
        int numBytes;

        close(fd[1]);

        if (read (fd[0],buffer,BUFFER_LEN) > 0)
        {
            printf("Child received \"%s\"\n",buffer);
        }

        close(fd[0]);
    }
    else
    {
        close(fd[0]);
        printf("Parent sending \"%s\"\n",TEXT);
        write(fd[1],TEXT,sizeof(TEXT));
        close(fd[1]);
    }

    return(EXIT_SUCCESS);
}
```

Inter-process communication: pipes (3)

- `dup(int oldFd)`
 - Copy file descriptor 'oldFd' into lowest available position
 - Lowest position generally obtained by `close()`-ing it immediately beforehand . . .
 - . . . or use `dup2()`
- `dup2(int oldFd, int newFd)`
 - `close(newFd)` (if open) then copy 'oldFd' into 'newFd'
 - If 'oldFd' is invalid, then 'newFd' is not closed
 - if 'oldFd==newFd', then nothing happens

Your turn!

Write a program that

- ₁ makes a pipe
- ₂ forks() two child processes
- ₃ one child runs '/usr/l^s', it sends its output to the pipe
- ₄ the other child runs '/usr/bin/wc', and gets its input from the pipe
- ₅ output of wc sent to stdout

Remember! close() unneeded file descriptors!

Inter-process communication: sockets (1)

- `socketpair()`

- Makes a pair of file descriptors (like `pipe()`)
- Two-way communication

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int socketpair(int domain, int type, int protocol, int
sv[ 2] );
```

- `domain` should be `AF_UNIX` (= `AF_LOCAL`)
 - `type` could be `SOCK_STREAM` (for TCP), `SOCK_DGRAM` (for UDP), `SOCK_SEQPACKET` (like TCP)
 - `protocol` can be 0 if there is only one choice for the given `type`.
 - `sv[]` receives pair of socket file descriptors, `sv[0]` and `sv[1]` indistinguishable
- Remember: `close()` unneeded file descriptor.

Inter-process communication: sockets (2)

```
/*
 * From ibm.com downloaded 2019-03-31
 */

/* This program fragment creates a pair of connected sockets
then
 * forks and communicates over them. Socket pairs have a two-
way
 * communication path. Messages can be sent in both directions.
 */

#include <stdlib.h>
#include <stdio.h>
#include <sys/socket.h>
#include <sys/types.h>
#define DATA1 "In Xanadu, did Kublai Khan..."
#define DATA2 "A stately pleasure dome decree..."

int main()
{
    int sockets[2], child;
    char buf[1024];
    if (socketpair(AF_UNIX, SOCK_STREAM, 0, sockets) < 0) {
        perror("opening stream socket pair");
        exit(EXIT_FAILURE);
    }
}
```

```
if ((child = fork()) == -1)
    perror("fork");
else if (child) { /* This is the parent. */
    close(sockets[0]);
    if (read(sockets[1], buf, 1024, 0) < 0)
        perror("reading stream message");
    printf("-->%s\n", buf);
    if (write(sockets[1], DATA2, sizeof(DATA2)) < 0)
        perror("writing stream message");
    close(sockets[1]);
} else { /* This is the child. */
    close(sockets[1]);
    if (write(sockets[0], DATA1, sizeof(DATA1)) < 0)
        perror("writing stream message");
    if (read(sockets[0], buf, 1024, 0) < 0)
        perror("reading stream message");
    printf("-->%s\n", buf);
    close(sockets[0]);
}
}
```


Your turn!

Write a program that:

- 1 creates a socketpair
- 2 `fork()`s a child process
- 3 the parent process gets text from user, send to socket
- 4 the child process gets text from socket, uppercases it, sends back to socket
- 5 the parent process gets text from socket and prints it

Inter-process communication: sockets (2)

- **Question:** *“Hey! I thought sockets were for communication between different machines!”*
- Server-side:
 - socket()
 - bind()
 - listen()
 - accept()
- Client-side:
 - getaddrinfo()
 - socket()
 - connect()
- ***Stay tuned!***



Inter-process communication: Message Queues (1)

- Asynchronous communication between processes
- Send “structs” between process
 - A little higher level than pipes/sockets, which send bytes
- `int msgget (int key, int flags)`
 - `key` is like a filename
 - `flags` is like for `open()` for files:
 - `(0660|IPC_CREAT)` to create
 - `0` to attach to existing queue
 - return msg queue id (like a file descriptor)

Say!

Where should that key come from?

- It is meant to come from `ftok()`:
 - `key_t ftok(const char *pathname, int proj_id);`
- Idea:
 - This user has access to files other users don't.
 - Therefore, build key based upon a file this user can access
- `pathname`
 - An existing, accessible file
- `proj_id`
 - A non-zero integer, whose lowest byte will be used
- (We will just enter an integer.)

Inter-process communication: Message Queues (2)

- `msgctl(int msgQId, int cmd, struct msqid_ds *buf)`
 - `msgQId`: which msg queue
 - `cmd`:
 - `IPC_STAT`: get info
 - `IPC_SET`: set info
 - `IPC_RMID`: delete queue
 - `buf`:
 - Information about the queue
 - `buf.msg_qbytes`: How big the queue should be

Inter-process communication: Message Queues (3)

- `msgsnd(int msgQId, const void* msgp, size_t msgsz, long msgtyp, int msgflg)`
 - `msgQId`: which queue
 - `msgp`: ptr to message
 - A struct with a `long` member var as first member
 - ***long value must be positive!***
 - `msgsz`: size of message
 - Does not count beginning long var
 - `msgtyp`:
 - “For whom”: same queue can be used by multiple readers.
 - ‘`msgtyp`’ tells for whom msg is for
 - `msgflg`: flags
 - Return value
 - 0 on success
 - -1 on failure
- `msgrcv(int msgQId, void* msgp, size_t msgsz, long msgtyp, int msgflg)`
 - `msgQId`: which queue
 - `msgp`: ptr to message
 - `msgsz`: size of message (not including long var)
 - `msgtyp`:
 - “For whom”: same queue can be used by multiple readers.
 - ‘`msgtyp`’ tells for whom msg is for
 - If 0 then first message of any type
 - `msgflg`: flags
 - Return value
 - number of bytes received
 - -1 on failure

Inter-process communication: Message Queues (4)

```
// msg.h
#define TEXT_LEN 80

struct AMessage
{
    long    msgType_;
    float   floatPt_;
    int     integer_;
    char    text_[TEXT_LEN];
};
```

```
// makeAndWriteMsgQueue.c

#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/msg.h>

#include "msg.h"

intgetMsgKey ()
{
    char text[TEXT_LEN];
    int toReturn;

    printf("Please enter a msg id integer: ");
    fgets(text,TEXT_LEN,stdin);
    toReturn = strtol(text,NULL,0);
    return(toReturn);
}
```

Inter-process communication: Message Queues (5)

```
int makeMsgQueue ()
{
    int toReturn;
    struct msqid_ds msgQInfoBuffer;

    do
    {
        int key = getMsgKey();

        toReturn = msgget(key,0660 | IPC_CREAT);
    }
    while (toReturn < 0);

    msgctl(toReturn,IPC_STAT,&msgQInfoBuffer);
    msgQInfoBuffer.msg_qbytes = 4096;
    msgctl(toReturn,IPC_SET,&msgQInfoBuffer);

    return(toReturn);
}
```

```
void writeMsgs (int msgQId)
{
    struct AMessage aMsg;
    char format[TEXT_LEN];
    char line[TEXT_LEN * 2];

    aMsg.msgType_ = 1;
    snprintf(format,TEXT_LEN,"%%f %%d %%%ds",TEXT_LEN);

    while (1)
    {
        aMsg.floatPt_ = 0.0;
        aMsg.integer_ = 0;
        memset(aMsg.text_,'\0',TEXT_LEN);

        printf("Please enter a float, int and word (or blank line to quit): ");
        fgets(line,TEXT_LEN*2,stdin);

        if (sscanf(line,format,&aMsg.floatPt_,&aMsg.integer_,aMsg.text_) <= 0)
        {
            msgsnd(msgQId,&aMsg,sizeof(aMsg)-sizeof(long),0);
            break;
        }

        aMsg.text_[TEXT_LEN-1] = '\0';
        msgsnd(msgQId,&aMsg,sizeof(aMsg)-sizeof(long),0);
    }
}
```

Inter-process communication: Message Queues (6)

```
int main ()
{
    int msgQId = makeMsgQueue();
    writeMsgs(msgQId);
    return(EXIT_SUCCESS);
}
```

```
// readAndDelMsgQueue.c
```

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/msg.h>
```

```
#include "msg.h"
```

```
int getMsgKey ()
{
    char text[TEXT_LEN];
    int toReturn;

    printf("Please enter a msg id integer: ");
    fgets(text, TEXT_LEN, stdin);
    toReturn = strtol(text, NULL, 0);
    return(toReturn);
}
```

Inter-process communication: Message Queues (7)

```
int main()
{
    struct AMessage msg;
    int msgKey = getMsgKey();
    int msgQId = msgget(msgKey, 0);

    if (msgQId < 0)
    {
        fprintf(stderr, "Sorry!\n");
        exit(EXIT_FAILURE);
    }

    while (1)
    {
        if (msgrcv(msgQId, &msg, sizeof(msg), 1, 0) < 0)
            break;

        if ((msg.floatPt_ == 0.0) && (msg.integer_ == 0) && (msg.text_[0] ==
'\0'))
            break;

        printf("Float:\t%g\nInt:\t%d\nWord:\t%s\n\n",
msg.floatPt_, msg.integer_, msg.text_
);
    }

    msgctl(msgQId, IPC_RMID, NULL);
    return(EXIT_SUCCESS);
}
```

\$ ipcs

- Command line tool for seeing active message queues (and shared memory, and semaphores)

\$ ipcrm -Q <key>

- Delete message queue by its key

Your turn!

Write an application:

Sender generates 2
random ints, sets
calculation_:

- QUIT_CALC
- ADD_CALC
- MULT_CALC

Receiver receives
message and either adds
or multiplies or quits

```
typedef enum
{
    QUIT_CALC,
    ADD_CALC,
    MULT_CALC
}
calc_ty;

struct MathProblem
{
    long msgType_;
    calc_ty calculation_;
    int int0_;
    int int1_;
};
```


Inter-process communication: shared memory (1)

- Allows multiple processes to access ***exact*** same pages
- If more than one process can write then should protect with semaphores
 - Much like `pthread_mutex_t` for threads
- Less of a use for this because threads already share memory

Inter-process communication: shared memory (2)

`void* mmap (void *addr, size_t len, int prot, int flags, int fd, off_t offset)`

- Memory map, does several things including making shared memory
 - `addr`: where to make shared (if `NULL` then OS will choose)
 - `len`: number of bytes
 - `offset`: starting byte (if '`fd`' is a file)
 - `fd`: file descriptor to map into memory
 - `prot`:
 - `PROT_EXEC`: may be executed
 - `PROC_READ`: may be read
 - `PROC_WRITE`: may be written
 - `PROC_NONE`: may be not be accessed
 - `flags`:
 - `MAP_SHARED`: visible to other processes

`int munmap (void* addr, size_t length)`

- Gets rid of memory map

Inter-process communication: shared memory (3)

```
// From a user named "slezica"
// From https://stackoverflow.com/questions/5656530/how-to-
// use-shared-memory-with-linux-in-c
// Downloaded 2019 Feb 3
#include <stdio.h>
#include <stdlib.h>
#include <sys/mman.h>
#include <string.h>
#include <unistd.h>

void* create_shared_memory(size_t size) {
    // Our memory buffer will be readable and writable:
    int protection = PROT_READ | PROT_WRITE;

    // The buffer will be shared (meaning other processes can
    // access it), but
    // anonymous (meaning third-party processes cannot obtain
    // an address for it),
    // so only this process and its children will be able to use it:
    int visibility = MAP_ANONYMOUS | MAP_SHARED;

    // The remaining parameters to `mmap()` are not important for
    // this use case,
    // but the manpage for `mmap` explains their purpose.
    return mmap(NULL, size, protection, visibility, 0, 0);
}
```

```
int main() {
    char* parent_message = "hello"; // parent process will write
    // this message
    char* child_message = "goodbye"; // child process will then
    // write this one

    void* shmem = create_shared_memory(128);

    int pid = fork();

    memcpy(shmem, parent_message,
    sizeof(parent_message));

    if (pid == 0) {
        sleep(1);
        printf("Child read: %s\n", shmem);
        memcpy(shmem, child_message, sizeof(child_message));
        printf("Child wrote: %s\n", shmem);
    } else {
        printf("Parent read: %s\n", shmem);
        sleep(2);
        printf("After 1s, parent read: %s\n", shmem);
    }
}
```

Inter-process communication: shared memory (4)

- Two other interfaces if want finer detail over which processes can access:
- Access with a “filename”
 - `shm_unlink()`, `shm_open()`
- Access with a key integer
 - `shmget()`, `shmctl()`, `shmat()`, `shmdt()`

Inter-process communication: shared memory (5)

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>
#include <unistd.h>
#include <errno.h>

#define SHARED_MEM_PATH "/TestSMem"

int main()
{
    const char* parent_message = "hello";
    const char* child_message = "goodbye";

    int fd;
    size_t size;
    void* addr;

    shm_unlink(SHARED_MEM_PATH);

    fd = shm_open(SHARED_MEM_PATH, O_CREAT|O_EXCL|O_RDWR, S_IRREAD|
S_IWRITE|S_IRGRP);
    size = 128;
    ftruncate(fd, size);

    addr = mmap(NULL, size, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);

    memcpy(addr, parent_message, sizeof(parent_message));
    printf("Original contents: %s\n", (char*)addr);

    int pid = fork();
```

```
    if (pid == 0) {
        printf("Child closing %d\n", fd);
        close(fd);
        printf("munmap()\n");
        munmap(addr, size);

        printf("shm_open()\n");
        errno = 0;
        fd = shm_open(SHARED_MEM_PATH, O_RDONLY, S_IRUSR|
S_IWUSR|S_IRGRP|S_IWGRP);
        size = 128;
        printf("fd = %d %s\n", fd, strerror(errno));


        addr = mmap(NULL, size, PROT_READ, MAP_SHARED, fd, 0);

        printf("Child read: %s\n", addr);

        printf("Child will attempt to write, but fail:\n");
        memcpy(addr, child_message, sizeof(child_message));
        printf("Child tried to write: %s\n", addr);

    } else {
        printf("Parent read: %s\n", addr);
        sleep(10);
        printf("After 10 seconds, parent read: %s\n", addr);
        shm_unlink(SHARED_MEM_PATH);
    }
}
```


Inter-process communication: shared memory (6)



Clever student's question *"Uh-oh! If two processes can write to the same memory at the same time, then we have the same problem as two threads writing to the same memory at the same time!"*

Semaphores (1)

- **Answer:** *“True! And we will use a similar solution”*
- **semaphore:** non-negative integer synchronization variable.
 - `sem_wait(s): { while(s==0){pause();}`
`s--; }`
 - `sem_post(s): { s++; }`
- OS guarantees that operations between brackets [] are executed indivisibly.
 - Only one `sem_wait()` or `sem_post()` operation at a time can modify `s`.
 - When `while` loop in `sem_wait()` terminates, only that `sem_wait()` can decrement `s`.
- **Semaphore invariant:** $(s \geq 0)$



Semaphores (2)

- What to include:
 - `#include <semaphore.h>`
- Types and functions:
 - `sem_t semaphore;`
 - `sem_init(sem_t* semPtr, int flag, int value)`
 - Initialize pointed-to semaphore, with `value`
 - *if **flag** == 1 then semaphore can be forked*
 - `sem_destroy(sem_t* semPtr)`
 - Destroy pointed-to semaphore. If it's negative then block.


Semaphores (3)

- `sem_wait(sem_t* semPtr)`
 - Decrement pointed-to semaphore. If it's negative then block.
- `sem_post(sem_t* semPtr)`
 - Increment pointed-to semaphore. Wake one blocked process if any.
- `sem_getvalue(sem_t* semPtr, int* valuePtr)`
 - Get value of pointed to semaphore.

Semaphores (4)

Question “*Where shall we put the semaphore so both processes can see it?*”

Answer “*Hey, didn’t we just make memory that both processes can see?*”



Semaphores in Shared Memory (1)

```
// semaEx.h
```

```
#define SHARED_MEM_PATH \  
"/TestSMem"
```

```
#define QUIT_CMD \  
"QUIT\n"
```

```
const int TEXT_LEN = 63;
```

```
struct SharedSegment  
{  
    sem_t sem_;  
    char isDataAvailable_;  
    char text_[TEXT_LEN];  
};
```

```
// Compile with  
// $ g++ shareSemaProd.cpp -o shareSemaProd -lpthread -lrt
```

```
#include <stdio.h>  
#include <stdlib.h>  
#include <sys/mman.h>  
#include <sys/stat.h>  
#include <fcntl.h>  
#include <string.h>  
#include <semaphore.h>  
#include <unistd.h>  
#include <errno.h>
```

```
#include "semaEx.h"
```

```
void writeToSegment(SharedSegment* segPtr)  
{  
    char quit[sizeof(QUIT_CMD)];  
    char* cPtr;  
  
    strcpy(quit, QUIT_CMD);  
    cPtr = strchr(quit, '\n');  
  
    if (cPtr != NULL)  
        *cPtr = '\0';  
  
    printf("Enter text (\"%s\" to quit): ", quit);  
    fgets(segPtr->text_, TEXT_LEN, stdin);  
    segPtr->isDataAvailable_ = 1;  
}
```

Semaphores in Shared Memory (2)

```
void      waitForMyTurn (SharedSegment* segPtr)
{
    int  isItMyTurn    = 0;

    while (1)
    {
        sem_wait(&segPtr->sem_);
        isItMyTurn = (segPtr->isDataAvailable_ == 0);
        sem_post(&segPtr->sem_);

        if (isItMyTurn)
            break;
        else
            sleep(1);
    }
}

int      main      ()
{
    SharedSegment* segmentPtr;

    shm_unlink(SHARED_MEM_PATH);

    int  fd = shm_open
        (SHARED_MEM_PATH,
         O_CREAT|O_EXCL|O_RDWR,
         S_IRREAD|S_IWWRITE|S_IRGRP|S_IWGRP
        );

    ftruncate(fd, sizeof(SharedSegment));
```

```
    segmentPtr = (SharedSegment*)mmap(NULL,
    sizeof(SharedSegment), PROT_READ | PROT_WRITE,
    MAP_SHARED, fd, 0);

    sem_init(&segmentPtr->sem_,1,1);
    segmentPtr->isDataAvailable_ = 0;

    do
    {
        writeToSegment(segmentPtr);
        waitForMyTurn(segmentPtr);
        printf("%s\n\n",segmentPtr->text_);
    }
    while ( strcmp(segmentPtr->text_,
    QUIT_CMD,sizeof(QUIT_CMD)-1) != 0);

    shm_unlink(SHARED_MEM_PATH);
    return(EXIT_SUCCESS);
}
```

Semaphores in Shared Memory (3)

```
// Compile with:  
// $ g++ shareSemaCons.cpp -o shareSemaCons -lpthread -lrt
```

```
#include <stdio.h>  
#include <stdlib.h>  
#include <sys/mman.h>  
#include <sys/stat.h>  
#include <fcntl.h>  
#include <string.h>  
#include <semaphore.h>  
#include <unistd.h>  
#include <errno.h>  
#include <ctype.h>
```

```
#include "semaEx.h"
```

```
void writeIntoSegment(SharedSegment* segPtr)  
{  
    printf("Received: %s",segPtr->text_);  
  
    for (char* cPtr = segPtr->text_; *cPtr != '\0'; cPtr++)  
        *cPtr = toupper(*cPtr);  
  
    printf("Sending: %s",segPtr->text_);  
}
```

```
void waitForMyTurn (SharedSegment* segPtr)  
{  
    int isItMyTurn = 0;  
  
    while (1)  
    {  
        sem_wait(&segPtr->sem_);  
        isItMyTurn = (segPtr->isDataAvailable_);  
        sem_post(&segPtr->sem_);  
  
        if (isItMyTurn)  
            break;  
        else  
        {  
            printf("Waiting\n");  
            sleep(1);  
        }  
    }  
}
```

```
int main()  
{  
  
    errno = 0;  
    int fd = shm_open  
              (SHARED_MEM_PATH,  
               O_RDWR,  
               S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP  
              );
```

Semaphores in Shared Memory (4)

```
SharedSegment*segPtr = (SharedSegment*)
    mmap(NULL,
        sizeof(SharedSegment),
        PROT_READ|PROT_WRITE,
        MAP_SHARED, fd, 0);
```

```
int  shouldQuit;
```

```
do
{
    waitForMyTurn(segPtr);
    writeIntoSegment(segPtr);
    shouldQuit = (strncmp(segPtr-
>text_,QUIT_CMD,sizeof(QUIT_CMD)-1) == 0);
    segPtr->isDataAvailable_ = 0;
}
while (!shouldQuit);


return(EXIT_SUCCESS);
}
```

Your turn!

Revise the last 2 programs to send math problems (like the message passing problem), but to send the answer back (like the socket pair problem).

Decisions! Decisions!

Which **shall** I use?



- Remember! These are only applicable for multi-process
 - Moot for multi-thread
 - They see the same memory
- One way?
 - Pipes!
- Two way?
 - Socket pairs!
 - **Bonus!** Straight-forward to generalize to ordinary multi-machine sockets
- LOTS OF DATA!
 - Shared memory
- What about messages?
 - Personally, I don't see that much of an advantage over pipes
 - If use pipes, revising to sockets is straight-forward

References:

- M. Tim Jones “*GNU/Linux Application Programming, 2nd Ed*” Course Technology, Cengage Learning. 2008
- ibm.com
- User “slezica” at stackoverflow.com