

Distributed Systems

Lecture 3

Networking Basics

Joseph Phillips

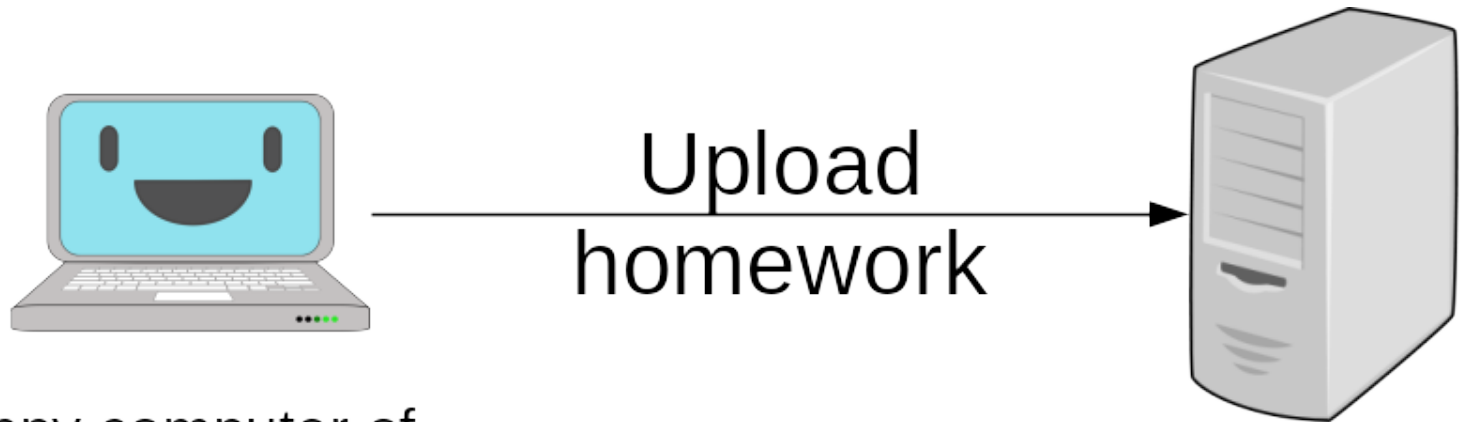
Copyright (c) 2019

Last modified 2019 September 12

Topics

- Networking
 - Ethernet
 - IP
 - TCP and UDP
- Byte ordering
- Character sets
 - Unicode
 - UTF 16
 - UTF 8
- Error handling
 - Java
 - C
 - C++
- Networking in Alternative Universes
 - Transport Layer Interface
 - Plan 9/Inferno

Motivation



Happy computer of
happy student who just
finished assignment

DePaul D2L
server

A happy student!

They just finished their assignment.

Now, if there only was a way to upload it?

Attempt #1



The student attaches computer to same network as D2L server

Attempt #1 (cont'd)

- **Question:** How can we get all computers to talk to each other on the same network?
- **Answer:** There are several approaches. Among the most common are:
 - Ethernet
 - IEEE 802.2/802.3
 - Wi-Fi
 - ZigBee


But first, make sure each ***network card*** is unique

- MAC: To distinguish one ***network interface*** from another
 - Media Access Control
 - 48-bits
 - Assigned by IEEE
 - network card manufacturers buy them in blocks
 - Unique for each ***network card***
 - Replace computer's network card? ***It has a new MAC address!***

Your turn!

What type of “computer”
would have more than one MAC address?

And now, an Ethernet Frame




- MAC addresses
 - **to** and **from** whom?
- Type: *What the heck am I carrying?*
 - 0x0800: IPv4
 - 0x86DD: IPv6
 - 0x0806: Address Resolution Protocol (ARP)
 - 0x8035: Reverse Address Resolution Protocol (RARP)
 - (There are many more)
- Data
 - An IP frame (stay tuned!)
- CRC
 - “Cyclic Redundancy Check”
 - Fixed-sized hash value
 - Good at detecting “burst errors” (sequence of wrong data)
 - Easy to compute in hardware

Variations on basic Ethernet


- IEEE 802.2/802.3
 - Similar to Ethernet frame
 - Steal some 8 bytes from data for more meta-info
- IEEE 802.11 (Wi-Fi), Has many more fields
 - Power management,
 - “this is a re-try frame”, etc.
- IEEE 802.15.4 (ZigBee): For Internet-of-Things
 - Multi-hop relay
 - Low power (e.g. battery)
 - Discovery of the number of devices in the network
 - Security

Ruh-roh!

Astute student:
*“DePaul is **too cheap**
to lay cable to every
student’s house.
Students are **too lazy**
to walk to DePaul!”*



Attempt #2




- Let's ***not*** use just ***one*** network.
- Let's use a ***network-of-networks!***
 - a.k.a. an ***internet***
 - “***i***nternet”: (small “i”) a network-of-networks
 - “***I***nternet”: (big “I”) ***The*** worldwide network-of-networks-of-networks

Attempt #2

Astute student: “*Hey!
All those networks have
different protocols!*”

*How can we get **any
machine** on one
network to talk to **any
machine** on another
network?”*



But first: Addressing computers on different networks


- With ***one network*** we could use the MAC address
- With ***an internet*** we need another scheme
- The ***IP address!***
 - Abstracts away from MAC address
 - A distributed registry of IP addresses to networks that host them (and know their MAC addresses)
- ***DNS*** (Domain Name Service)
 - Computers love integers, but humans love strings
 - IP address => name
 - name => IP address

And now . . . Attempt #2: Talking

- We can address another computer on another network
- Now we have to ***talk***
- There really are two issues here . . .
 - A multi-network issue (“hardware”)
 - A multi-protocol issue (“software”)

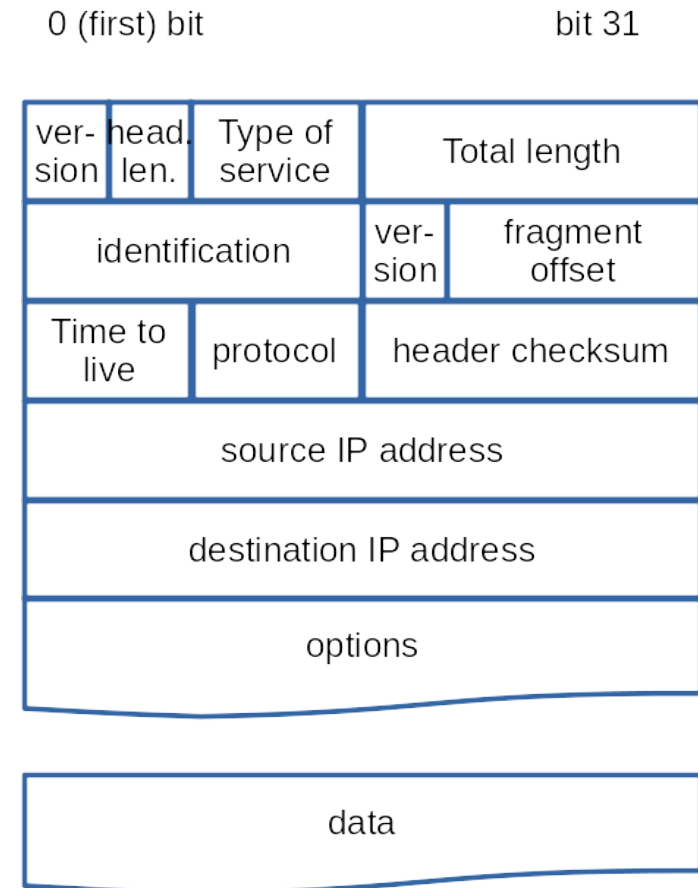
Attempt #2: “Hardware”

- Introducing the **bridge**
 - a device to bridge different networks
- A **router**
 - bridge between different **types** of networks
 - e.g. Ethernet and Asymmetric Digital Subscriber Line (ADSL)
 - e.g. Your home wireless



Attempt #2: “Software”

- IP
 - Says where it is going with an **IP address**
 - Not just a **this-network-only** MAC address
 - Version
 - E.g. IPv4
 - Type of service:
 - minimize delay, or,
 - maximize throughput, or,
 - maximize reliability, or,
 - minimize monetary cost
 - Identification
 - Integer incremented for subsequent datagrams
 - Time to live:
 - How many more “hops” it can do before it dies
 - Prevents infinite looping



IPv4 vs. IPv6

	IPv4	IPv6
Address	4 octets (32 bits)	16 octets (128 bits)
Notation	Decimal-dot	Hexadecimal-colon
Security	Achievable by higher levels	Encryption and authentication at this level
Maximum packet	65535 bytes	4,294,867,295 bytes ("jumbogram")

But wait! There's more . . .

- Besides IP, there are other low-level protocols
 - IGMP: Internet Group Management Protocol
 - Manages multi-cast groups
 - ICMP: Internet Control Message Protocol
 - Router-to-router communication about errors, etc.
 - Used by ping, traceroute

ping

- “Are you there? / Can I reach you?”

\$ **ping** *www.google.com*

PING www.google.com (172.217.4.36) 56(84) bytes of data.

64 bytes from lga15s46-in-f4.1e100.net (172.217.4.36): icmp_seq=1 ttl=54
time=11.4 ms

64 bytes from lga15s46-in-f4.1e100.net (172.217.4.36): icmp_seq=2 ttl=54
time=11.6 ms

64 bytes from lga15s46-in-f4.1e100.net (172.217.4.36): icmp_seq=3 ttl=54
time=9.11 ms

64 bytes from lga15s46-in-f4.1e100.net (172.217.4.36): icmp_seq=4 ttl=54
time=9.15 ms

64 bytes from lga15s46-in-f4.1e100.net (172.217.4.36): icmp_seq=5 ttl=54
time=8.42 ms

^C

--- www.google.com ping statistics ---

5 packets transmitted, 5 received, 0% packet loss, time 4006ms

rtt min/avg/max/mdev = 8.422/9.960/11.630/1.335 ms

traceroute

- “How did my packet get to you?”


\$ **traceroute 140.192.36.187**

traceroute to 140.192.36.187 (140.192.36.187), 30 hops max, 60 byte packets

```
1 gateway (192.168.0.1) 1.807 ms 2.609 ms 3.661 ms
2 96.120.24.253 (96.120.24.253) 21.038 ms 21.320 ms 21.968 ms
3 be141-sur07.area4.il.chicago.comcast.net (68.85.179.161) 22.698 ms 22.336 ms 24.710 ms
4 be-110-ar01.area4.il.chicago.comcast.net (68.86.184.253) 27.049 ms 27.665 ms 27.357 ms
5 be-33491-cr02.350ecermak.il.ibone.comcast.net (68.86.91.165) 28.351 ms 31.459 ms
  31.797 ms
6 be-10563-pe01.350ecermak.il.ibone.comcast.net (68.86.82.158) 29.445 ms 12.555 ms
  21.235 ms
7 50.242.150.34 (50.242.150.34) 22.936 ms 22.140 ms 22.447 ms
8 0.ae4.cr2.ord6.scnnet.net (204.93.204.87) 23.143 ms 23.439 ms 0.ae12.cr2.ord6.scnnet.net
  (204.93.204.83) 25.377 ms
9 0.ae2.ar10.ord6.scnnet.net (204.93.204.81) 24.552 ms 24.166 ms 25.500 ms
10 unknown.servercentral.net (50.31.167.142) 31.764 ms 31.978 ms 31.264 ms
11 140.192.10.5 (140.192.10.5) 33.778 ms 34.763 ms 34.340 ms
12 mfc-cst-5e-te13-6.netequip.depaul.edu (140.192.9.229) 16.235 ms 20.583 ms 19.894 ms
```

Ruh-roh again!

- Astute student “*Hey, what you propose is an **unreliable stream of bytes**, it has **no meaning**.*”




Attempt #3


- Let us build a protocol on top of IP
- Nay, let us build 2 protocols
 - UDP
 - TCP

But first, make sure we are talking to the correct service

- Introducing the ***port***
 - Implies which service
 - An integer 1 to 65535 ***and*** a protocol (TCP or UDP)
 - Example:
 - ssh => TCP port 22; http => TCP port 80
 - Both ssh and http packets can travel same path, but not get confused



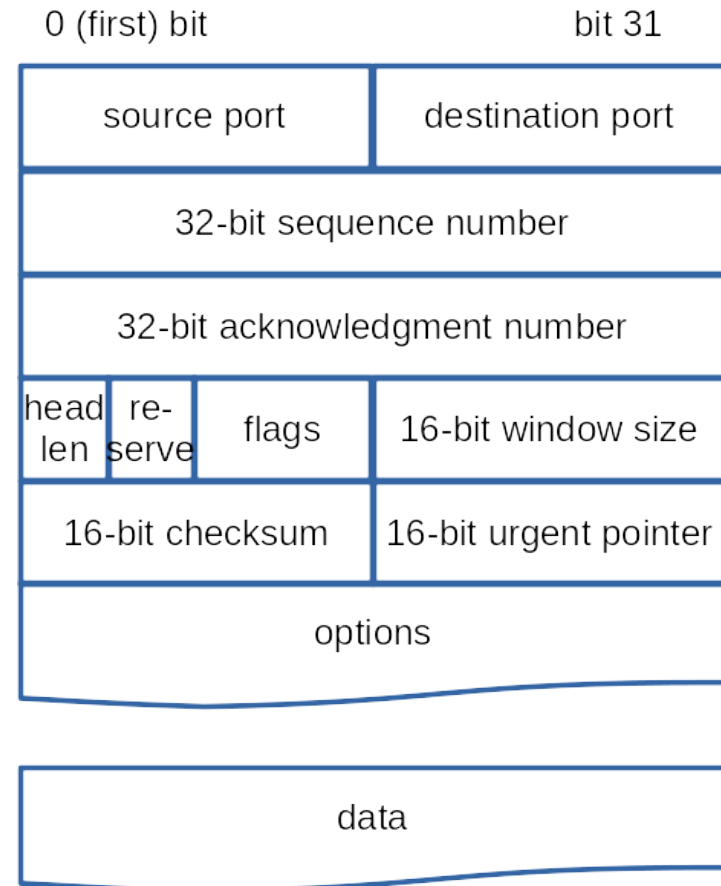
And now, UDP



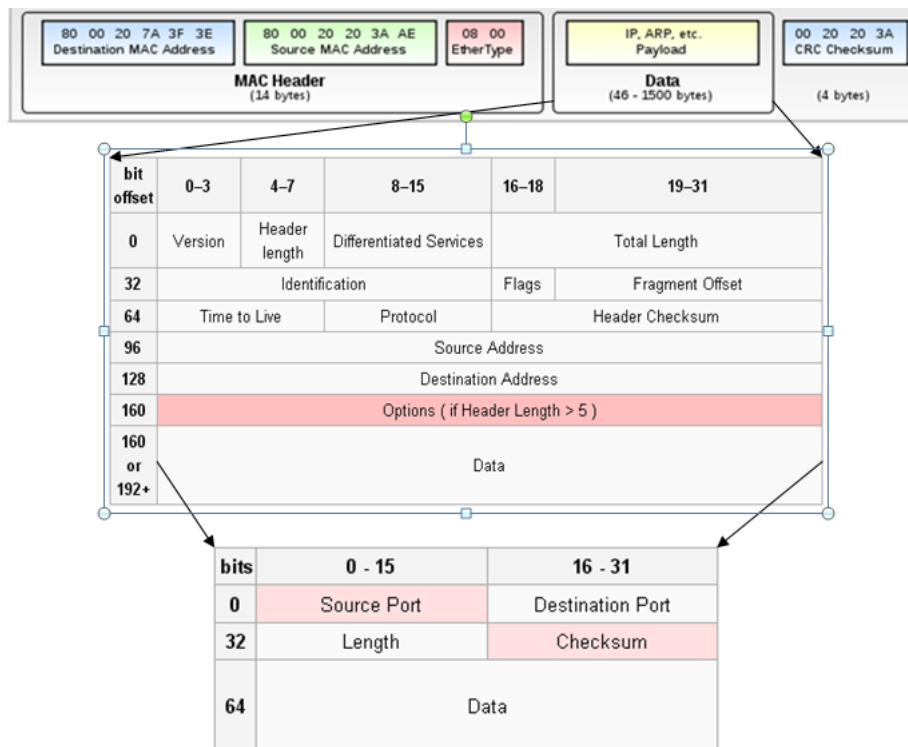
- Port
 - Implies which service on a given machine
 - TCP port 22: Secure SHell (ssh)
 - TCP port 23: Telnet
 - TCP port 80: http
 - TCP port 443: https
- Checksum
 - Helps detect errors

TCP

- Ports, checksum, and . . .
- Sequence number
 - For receiver to reassemble stream
- Acknowledgment number
 - Next sequence number expected



Yay! TCP/IP Saves the Day!



Application (e.g. FTP)

Transport (e.g. UDP)


Internet (e.g. IP)

Network (e.g. Ethernet)

Hardware (e.g. voltages)


Well, not quite: must consider application layer interoperability

I speak Big Endian



I kaeps elttil naidnE

I speak ASCII



SJIS で話します

Byte-ordering: Java

- We define the network ordering to be Big Endian
- Java coder? You're covered!
 - Java is defined to represent data in Big Endian

Byte-ordering: C/C++

- We define the network ordering to be Big Endian
- C/C++ coder? Just a *little* more work . . .

```
#include <arpa/inet.h>
```

```
uint32_t htonl(uint32_t hostlong);
```

```
uint16_t htons(uint16_t hostshort);
```

```
uint32_t ntohl(uint32_t netlong);
```

```
uint16_t ntohs(uint16_t netshort);
```

And if you *know* the endian you want

```
#include <endian.h>
```

```
uint16_t htobe16(uint16_t host_16bits); // Host To Big Endian  
uint16_t htole16(uint16_t host_16bits); // Host To Little Endian  
uint16_t be16toh(uint16_t big_endian_16bits);  
uint16_t le16toh(uint16_t little_endian_16bits);
```

```
uint32_t htobe32(uint32_t host_32bits);  
uint32_t htole32(uint32_t host_32bits);  
uint32_t be32toh(uint32_t big_endian_32bits);  
uint32_t le32toh(uint32_t little_endian_32bits);
```

```
uint64_t htobe64(uint64_t host_64bits);  
uint64_t htole64(uint64_t host_64bits);  
uint64_t be64toh(uint64_t big_endian_64bits);  
uint64_t le64toh(uint64_t little_endian_64bits);
```

Your Turn!

Write your own

```
uint64_t htonll(uint64_t hostLongLong)
```

```
uint64_t ntohll(uint64_t netLongLong)
```

Character Sets

- Unicode
 - Extends ASCII
 - Defines 1,112,064 “code points”
 - Chars of living and dead langs
 - control chars
- Represent in hexadecimal
 - U + *hhhh*
- Does ***not*** say how to represent bits of *hhhh*.



Unicode, cont'd

- UTF-32:
 - 4 bytes used to represent each char
 - Advantage:
 - Straight-forward
 - Disadvantage:
 - Wasteful of bytes
 - Used by:
 - C/C++ internal `wchar_t`
- UTF-16:
 - Either 2 bytes or 4 bytes
 - Advantage:
 - Less bytes than UTF-32 (for most chars but Chinese)
 - Disadvantage:
 - More complicated to handle 2 byte vs 4 byte chars
 - Used by:
 - Java internal `Char`

Unicode, cont'd

- UTF-8
 - Uses 1, 2, 3 or 4 “octets” (8-bit bytes)
 - Advantage:
 - More efficient than even UTF-16 (esp. for Latin chars)
 - Disadvantage:
 - More complicated to handle variable-lengthed chars
 - Used by:
 - Unix
 - Web pages

Unicode in C/C++

- Includes:
 - `wchar.h`, `locale.h`
- `locale`:
 - `LC_COLLATE`: string collation
 - `LC_CTYPE`: character type
 - `LC_MESSAGES`: natural language messages
 - `LC_MONETARY`: money formatting
 - `LC_NUMERIC`: number formatting
 - `LC_TIME`: date and time formatting
 - `LC_ALL`: Everything
 - `setlocale(LC_ALL, "");` // Set default for current system
 - `setlocale(LC_ALL, "en_US.UTF-8");` // Set for particular lang and charset
- Types:
 - `wchar_t`, `wchar_t*`
- Constants:
 - `wchar_t hiraganaA = L'\x3042';`
 - `wchar_t* L"日本語で";`

Unicode in C: Output wchar_t both internal & external

```
#include <stdio.h>
#include <locale.h>
#include <wchar.h>
```

```
int wprintf(const wchar_t *format, ...);
int fwprintf(FILE *stream, const wchar_t *format, ...);
int swprintf(wchar_t *wcs, size_t maxlen,
             const wchar_t *format, ...);
```

```
setlocale(LC_ALL,"en_US.UTF-8");
wprintf(L"char: '%c' wchar_t: '%lc'\n",'A',L'\x3042');
wprintf(L"char*: %s, wchar_t* %ls\n","in English",L" 日本語で");
```

Unicode in C: Output

wchar_t internal, char external

```
void    printUtf8Char (wchar_t  c
                        )
{

    if    (c <= 0x7F)
        putchar(c);
    else
        if    (c <= 0x7FF)
        {
            putchar( 0xC0 | (c >> 6) );
            putchar( 0x80 | (c & 0x3F) );
        }
    else
```

Unicode in C: Output

wchar_t internal, char external

```
if (c <= 0xFFFF)
{
    putchar( 0xE0 | (c >> 12) );
    putchar( 0x80 | (0x3F & (c >> 6)) );
    putchar( 0x80 | (0x3F & c) );
}
else if (c <= 0x1FFFFFF)
{
    putchar( 0xF0 | (c >> 18) );
    putchar( 0x80 | (0x3F & (c >> 12)) );
    putchar( 0x80 | (0x3F & (c >> 6)) );
    putchar( 0x80 | (0x3F & c) );
}
}
```

Your turn!

- Japanese has two alphabets:
 - Hiragana
 - Katakana
- For Hiragana:
 - The first char is あ (say “ah”) (U+3042)
 - The last char is ん (do not say, but pronounce “n”) (U+3093)
- Write a program to print the chars of Hiragana

Unicode in C: Input

- Includes:
 - `#include <wchar.h>`, `#include <stdlib.h>`
- `wchar_t *fgetws(wchar_t *ws, int n, FILE *stream);`
- Conversion from UTF-8/UTF-16 to `wchar_t`:
 - `mbstowcs()`, `mbsrtowcs()`
 - Multi-byte to wide char
 - `wcstombs()`, `wcrtomb()`
 - Wide char to multi-byte
 - Uses `LC_TYPE` of `locale()`

mbrtowc() example

```
// PURPOSE: To translate the unicode sequence pointed to by 'charPtr'  
// into a wchar_t C-string. Resulting wchar_t C-string placed in 'dest'  
// of length 'destLength', if the resulting string would fit. Otherwise,  
// space allocated on heap. In either case 'numWideChars' is set to the  
// number of wchar_t instances that result (not including the ending  
// '\0'). Returns resulting wchar_t C-string (either 'dest' or heap  
// address).
```

```
wchar_t*  utf8CPtrToWCPtr (const char*  charPtr,  
                           size_t&      numWideChars,  
                           wchar_t*      dest,  
                           size_t        destLength  
                           )
```

```
{  
    // I. Application validity check:  
    if ( (charPtr == NULL) || (dest == NULL) )  
        throw L"NULL ptr to charPtrToWcharPtr()";
```

```
    // II. Translate:  
    wchar_t*  space      = dest;  
    size_t    spaceLen    = destLength - 1;  
    wchar_t*  spaceEnd    = space + spaceLen;
```

```
    // II.A. Initialize 'mbstate':  
    mbstate_t  mbstate;  
    memset(&mbstate, '\0', sizeof(mbstate_t));
```

```
    // II.B. Each iteration translates to one 'wchar_t':  
    wchar_t*  spaceRun;  
    size_t    charPtrStep;  
    size_t    length      = strlen(charPtr);
```

```
    for ( spaceRun = space;  
          ;  
          spaceRun++, charPtr += charPtrStep, length -= charPtrStep  
          )  
    {  
        // II.B.1. Translate to wchar_t and see how many more chars to  
        // advance in 'charPtr':  
        wchar_t  resultWC;
```

```
        charPtrStep = mbrtowc(&resultWC, charPtr, length, &mbstate);
```

```
        // II.B.2. Stop loop if at end:  
        if (charPtrStep == 0)  
            break;
```

```
        // II.B.3. Handle errors:  
        if (charPtrStep == (size_t)-1)  
            throw L"illegal unicode byte sequence";
```

```
        // II.B.4. Allocate more space if needed:  
        if (spaceRun == spaceEnd)  
        {
```

```
            wchar_t*  tempWCPtr;  
            size_t      numWCharsToCopy = spaceRun - space;
```

```
            spaceLen    += spaceLen;  
            tempWCPtr    = (wchar_t*)malloc((spaceLen+1)*sizeof(wchar_t));
```

```
            if (tempWCPtr == NULL)  
                throw L"calloc() failure in charPtrToWcharPtr()";
```

mbrtowc() example, cont'd

```
memcpy(tempWCPtr,space,numWCharsToCopy*sizeof(
wchar_t));
```

```
    if (space != dest)
        free(space);
```

```
    space      = tempWCPtr;
    spaceRun    = space + numWCharsToCopy;
    spaceEnd    = space + spaceLen;
}
```

```
// II.B.5. Store char:
*spaceRun = resultWC;
}
```

```
// II.C. End wide char string:
*spaceRun = L'\0';
```

```
// IV. Finished:
numWideChars    = spaceRun - space;
return(space);
}
```

Error Handling: Java

- Exceptions (From Vaibhav Aggarwal
<http://vaibhavblogs.org/2012/12/common-java-networking-exceptions/>):
 - BindException
 - ClosedChannelException
 - ConnectException
 - InterruptedIOException
 - NoRouteToHostException
 - PortUnreachableException
 - ProtocolException
 - SocketException
 - SocketTimeoutException
 - SSLException
 - UnknownHostException
 - UnknownServiceException

Error handling: C

- include
 - errno.h, string.h
- Defines global var 'errno' that holds error code for last system call
 - Match with constants:
 - E2BIG Argument list too long
 - EACCES Permission denied
 - EADDRINUSE Address already in use
 - EADDRNOTAVAIL Address not available
 - EAFNOSUPPORT Address family not supported
 - EAGAIN Resource temporarily unavailable
 - There are many more!
 - char* strerror (int errnum)
 - Returns ptr to error message for 'errnum'
 - void perror(const char* s)
 - Prints 's', then the error message, to stderr

Your turn!

- What is wrong with the following:

```
if (somecall() == -1) {  
    printf("somecall() failed\n");  
    if (errno == ...) { ... }  
}
```

- How would you fix it?

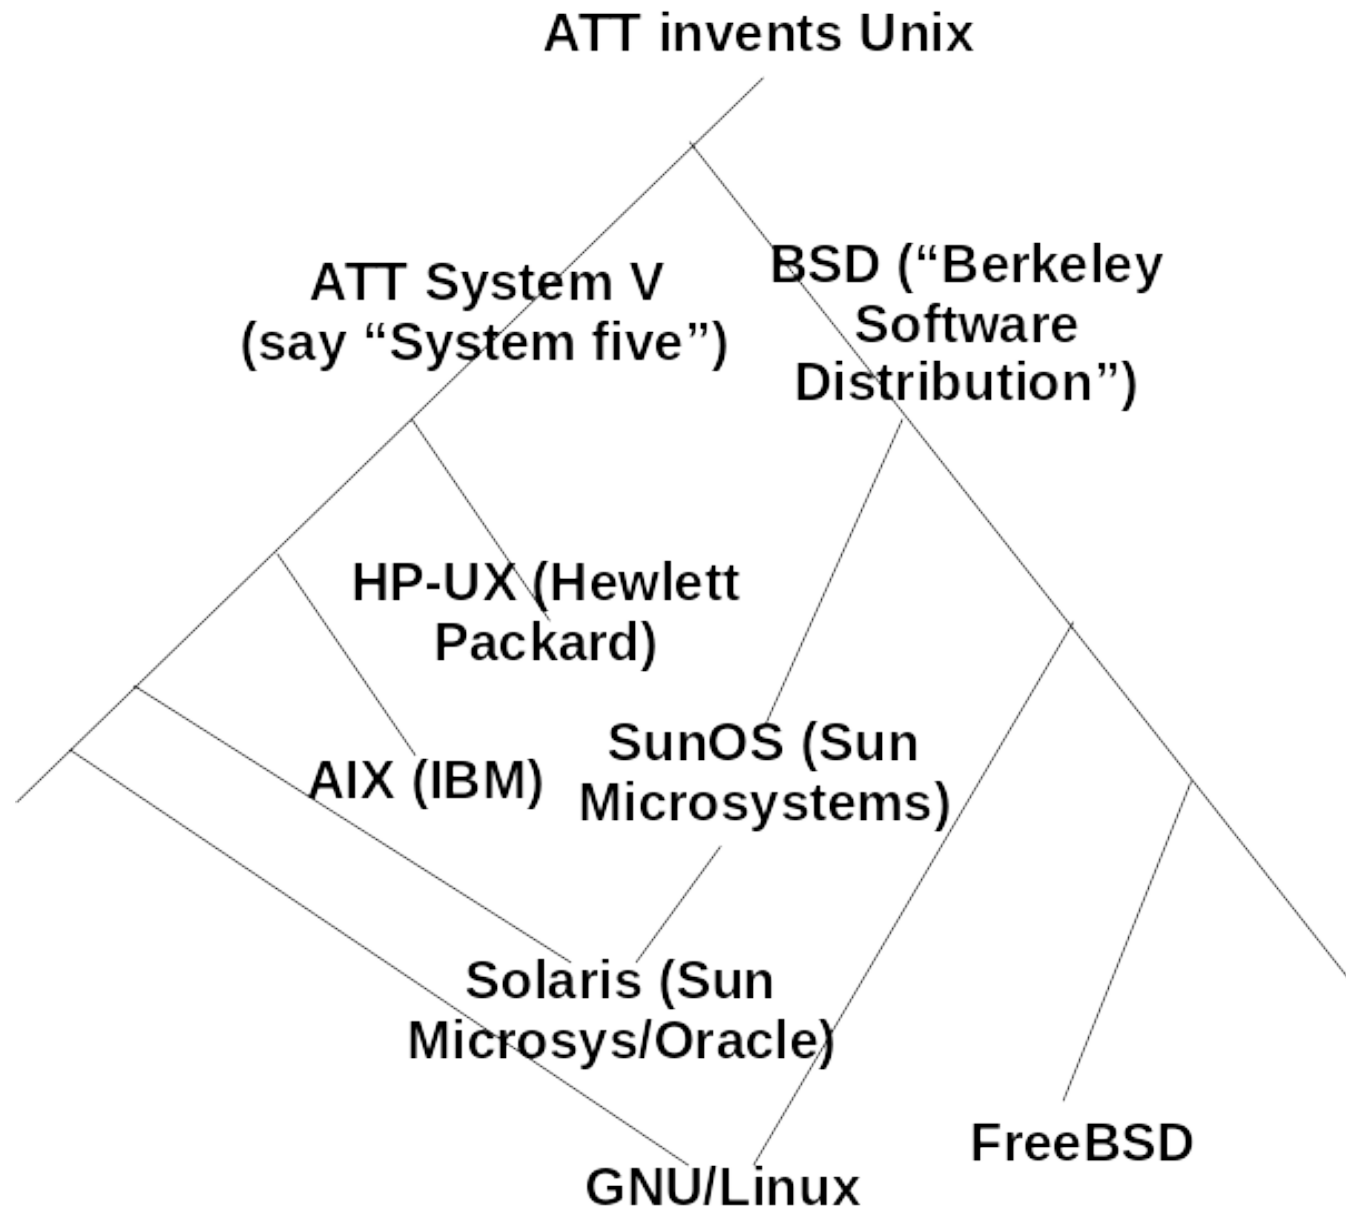
Did GOD ordain that
networking look like this?



Not quite . . .
In a galaxy far, far away



A History of Derivations of Unix




Alternatives to sockets

- Berkeley Software Distribution (BSD)
 - Sockets
- ATT System V (Pronounced “System five”)
 - “Transport Layer Interface”

BSD sockets	System V TLI
read, write, close	read, write, t_close
socket, socketpair	t_open
bind, listen	t_bind, t_unbind
getsockopt, setsockopt, select	t_optmgmt, t_look
strerror	t_strerror
connect	t_connect
accept	t_listen, t_accept, t_snddis

Plan 9/Inferno

- Distributed OS
- From Bell Labs
- Derived from Unix
- Big ideas of Plan 9/Inferno
 - 1 “Everything is a file” (even more so than Unix)
 - 2 Computable namespaces
 - Generalization of filepaths
 - 3 File service protocol
- Name from 1958 movie called “the worst ever made”, “*Plan 9 from Outer Space*”



Plan 9 from Bell Labs

“Everything is a file”

- Background, consider Linux’s /proc directory
 - presents information on processes
 - appears ***as if*** it were in files in directories
- But Plan 9/Inferno takes it further:
 - networking and other services used as if they were files

Echo server in Plan 9

```
// The Organization of Networks in Plan 9
// Dave Presotto, Phil Winterbottom
// presotto,philw@plan9.bell-labs.com
int echo_server(void)
{
    int dfd, lcfld;
    char adir[40], ldir[40];
    int n;
    char buf[256];

    int afd = announce("tcp!*echo", adir);

    if(afd < 0)
        return -1;

    for(;;)
    {
        /* listen for a call */
        lcfld = listen(adir, ldir);
        if (lcfld < 0)
            return -1;
    }
```

```
/* fork a process to echo */
switch (fork())
{
    case 0:
        /* accept the call and open the data file */
        dfd = accept(lcfld, ldir);
        if (dfd < 0)
            return -1;
        /* echo until EOF */
        while ((n = read(dfd, buf, sizeof(buf))) > 0)
            write(dfd, buf, n);
        exits(0);
    case -1:
        perror("forking");
    default:
        close(lcfld);
        break;
}
}
```

Plan 9/Inferno System calls

- Designed as a distributed form of Unix
 - exits()
 - Like C/Unix exit(), but returns a string
 - So error codes do not have to be standardized across systems
 - rfork(int flags)
 - Like C/Unix fork(), but more general
 - Flags include:
 - RFMEM If set, the kernel will force sharing of the entire address space, typically by sharing the hardware page table directly. (**QUESTION: What does this remind you of?**)
 - RFNOWAIT If set, the child process will be dissociated from the parent. Upon exit the child will not leave a status for the parent to collect. (**QUESTION: What does this remind you of?**)
 - RFFDG If set, the invoker's file descriptor table is copied; otherwise the two processes share a single table.

Plan 9/Inferno System calls, cont'd

- `void* rendezvous (void* tag, void* value)`
 - allows two processes to synchronize and exchange a value.
 - In conjunction with the shared memory system calls, it enables parallel programs to control their scheduling.

References:

- W. Richard Stevens “*TCP/IP Illustrated, Volume 1*” Addison-Wesley Professional Computing Series. 1994
- Douglas E. Comer “*Internetworking with TCP/IP: Principles, Protocols, and Architecture. 6th Ed.*” Pearson Education, 2014.
- Vaibhav Aggarwal <http://vaibhavblogs.org/2012/12/common-java-networking-exceptions/>
- Dave Presotto, Phil Winterbottom “*The Organization of Networks in Plan 9*”
- Charles Forsyth “*The Name Game: Featuring Plan 9 and Inferno*” YouTube talk
- http://man.cat-v.org/plan_9/2/rendezvous (man page for rendezvous() system call)
- Joseph Phillips, Applied Philosophy of Science