

Distributed Systems

Lecture 8

Server-side Sockets

Joseph Phillips

Copyright (c) 2019

Last modified 2019 October 29

Reading

- Chapter 9 “*Sockets for Servers*” of Elliotte Rusty Harold “*Java Network Programming: 4th Ed.*”

Topics

- Server-side sockets in Java
- Server-side sockets in C

Motivation

Last lecture we talked about implementing clients without the new-fangled Java tools of `URLConnection`, etc.

- Very nice
- But who is the client talking to?

Answer: “***The server!***”

Sockets!

- Remember . . .
- Duplex!
 - Can both read and write to same object
- Use ports
 - An integer from 1..65535
 - Acts like a “mailbox”



Java: the ServerSocket class

- Constructors:
 - `public ServerSocket(int port)` throws `BindException`, `IOException`
 - `public ServerSocket(int port, int queueLength)` throws `BindException`, `IOException`
 - `public ServerSocket(int port, int queueLen, InetAddress bindAddr)` throws `IOException`
 - `public ServerSocket()` throws `IOException`
- Accessors:
 - `public Socket accept()`
 - `public void bind (SocketAddress endpoint)` throws `IOException`
 - `public void bind (SocketAddress endpoint, int queueLen)` throws `IOException`

DaytimeServer.java

```
import java.net.*;
import java.io.*;
import java.util.Date;

public class DaytimeServer
{
    public final static int PORT = 1024;

    public static void main (String[] args)
    {
        System.out.println("Connect to port " +
PORT);
        try (ServerSocket server = new
ServerSocket(PORT))
        {
            while (true)
            {
                try (Socket connection = server.accept())
                {
                    Writer out;

                    out = new OutputStreamWriter
                        (connection.getOutputStream());
                    Date now = new Date();
                    out.write(now.toString() + "\r\n");
                    out.flush();
                    connection.close();
                }
                catch (IOException ex)
                {
                    {}
                }
            }
        }
        catch (IOException ex)
        {
            System.err.println(ex);
        }
    }
}
```

But wait!

This server is single-threaded!
(Slower than it needs to be)

Two approaches for basic threads in Java:

- Extend class Thread,
and @Override
public void
run().
- Then make a new
Thread of that class,
and have it do
start() method.
- Implement interface
Runnable (which also
needs @Override
public void run())
- Make a new Thread,
giving instance of the
new class as an
argument to the
constructor

MultithreadedDaytimeServer.java

```
import java.net.*;
import java.io.*;
import java.util.Date;

public class MultithreadedDaytimeServer
{
    public final static int PORT = 1024;

    public static void main (String[] args)
    {
        System.out.println("Connect to port " + PORT);
        try (ServerSocket server = new ServerSocket(PORT))
        {
            while (true)
            {
                try
                {
                    Socket connection = server.accept();
                    Thread task = new DaytimeThread(connection);
                    task.start();
                }
                catch (IOException ex)
                {}
            }
        }
    }
}
```

```
        catch (IOException ex)
        {
            System.err.println(ex);
        }
    }

    private static class DaytimeThread extends
Thread
    {
        private Socket connection;

        DaytimeThread(Socket connection)
        {
            this.connection = connection;
        }

        @Override
        public void run ()
        {
            try
            {
```

MultithreadedDaytimeServer.java

```
    Writer out = new OutputStreamWriter(connection.getOutputStream());
    Date now = new Date();
    out.write(now.toString() + "\r\n");
    out.flush();
}
catch (IOException ex)
{
    System.err.println(ex);
}
finally
{
    try
    {
        connection.close();
    }
    catch (IOException e)
    {
        // ignore
    }
}
}
```

Your turn!

From a security point of view:
what is the problem with spawning a new
thread every time someone connect()s?

More advanced Threads in Java

- Make a class that implements `Callable<WhateverClass>`

A better idea: Thread Pools!

```
import java.net.*;
import java.io.*;
import java.util.*;
import java.util.concurrent.*;

public class PooledDaytimeServer
{
    public final static int PORT = 1024;

    public static void main (String[] args)
    {
        System.out.println("Connect to port " + PORT);
        ExecutorService pool = Executors.newFixedThreadPool(16);

        try (ServerSocket server = new ServerSocket(PORT))
        {
            while (true)
            {
                try
                {
                    Socket connection = server.accept();
                    Callable<Void> task = new DaytimeTask(connection);

                    pool.submit(task);
                }
            }
        }
    }
}
```

```
        catch (IOException ex)
        {
        }
    }
}
catch (IOException ex)
{
    System.err.println(ex);
}
}

private static class DaytimeTask implements Callable<Void>
{
    private Socket connection;

    DaytimeTask(Socket connection)
    {
        this.connection = connection;
    }

    @Override
    public Void call ()
    {
        try
        {
        }
    }
}
```

A better idea: Thread Pools!

```
Writer out = new OutputStreamWriter(connection.getOutputStream());
Date now = new Date();
out.write(now.toString() + "\r\n");
out.flush();
}
catch (IOException ex)
{
    System.err.println(ex);
}
finally
{
    try
    {
        connection.close();
    }
    catch (IOException e)
    {
        // ignore
    }
}
return null;
}
}
```

Your turn!

Write an uppercasing server.

Client connects and gives line.

Server uppercases text
and sends back to client

More class Socket goodies

- `public void setSoTimeout (int milliseconds)`
throws `SocketException`
- `public int getSoTimeout ()` throws `SocketException`
 - How long to wait for `read()`
 - `milliseconds == 0` means “Wait forever”
 - After time expires throws `InterruptedException`
 - Prepare to catch it!
 - Socket still open, next `read()` might succeed

More class Socket goodies

- `public void setReuseAddress (boolean on)` throws `SocketException`
- `public boolean getReuseAddress ()` throws `SocketException`
 - Allow another socket to bind to same port immediately after `close()`?
 - false by default

More class Socket goodies

- `public void setReceiveBufferSize (int size)`
throws `SocketException`,
`IllegalArgumentException`
- `public int getReceiveBufferSize ()` throws
`SocketException`

More class Socket goodies

- `public void setPerformancePreferences (int connectionTime, int latency, int bandwidth)`
 - Sets relative preferences among connection-time, latency and bandwidth
 - `ss.setPerformancePreferences(2,1,3)`
 - Absolute values do no matter, only relative to each other
 - Example:
 - connection-time (=2) 2nd most important
 - latency (=1) least important
 - bandwidth (=3) most important
 - Implementation depends on particular Java VM

Server sockets in C

```
#include <sys/socket.h> //For socket()
#include <netinet/in.h> //For sockaddr_in and htons()
#include <netdb.h>       //For getaddrinfo()
#include <errno.h>       //For errno var
#include <sys/stat.h>    //For open(), read(), write()
#include <fcntl.h>       // and close()
```

- `socket()`: Ask OS for a socket
- `bind()`: Bind socket and port together
- `listen()`: Tell how many clients may queue
- `accept()`: Wait until a client connects
- `write()`: Write to client/server
- `read()`: Read from client/server
- `close()`: Close socket with client/server.

socket()

```
// Create a socket
int socketDescriptor =
    socket(AF_INET,      // AF_INET domain
          SOCK_STREAM,  // Reliable TCP
          0);
```

- Returns
 - A file descriptor that the server uses to see if a client has connected, or,
 - -1 on error
- There's also `SOCK_DGRAM` for UDP
- Last parameter type if used for `SOCK_RAW`

bind()

```
// Bind socket to port
// We'll fill in this datastruct
struct sockaddr_in socketInfo;

// Fill socketInfo with 0's
memset(&socketInfo, '\0', sizeof(socketInfo));
// Use std TCP/IP
socketInfo.sin_family = AF_INET;
// Tell port in network endian with htons()
socketInfo.sin_port = htons(portNumber);
    // (1) Allow connections from same machine only:
    struct in_addr addr;
    if (inet_aton("127.0.0.1", &addr) == 0) exit(EXIT_FAILURE);
    socketInfo.sin_addr.s_addr = addr.s_addr;
    // or (2) Allow machine to connect to this service
    socketInfo.sin_addr.s_addr = INADDR_ANY;
// Try to bind socket with port and other specifications
int status = bind(socketDescriptor, // from socket()
                  (struct sockaddr*)&socketInfo,
                  sizeof(socketInfo));
status == -1 on error
```

bind()

```
// Bind socket to port
// We'll fill in this datastruct
struct sockaddr_in socketInfo;

// Fill socketInfo with 0's
memset(&socketInfo, '\0', sizeof(socketInfo));
// Use std TCP/IP
socketInfo.sin_family = AF_INET;
// Tell port in network endian with htons()
socketInfo.sin_port = htons(portNumber);
    // (1) Allow connections from same machine only:
    struct in_addr addr;
    if (inet_aton("127.0.0.1", &addr) == 0) exit(EXIT_FAILURE);
    socketInfo.sin_addr.s_addr = addr.s_addr;
    // or (2) Allow machine to connect to this service
    socketInfo.sin_addr.s_addr = INADDR_ANY;
// Try to bind socket with port and other specifications
int status = bind(socketDescriptor, // from socket()
                  (struct sockaddr*)&socketInfo,
                  sizeof(socketInfo));
status == -1 on error
```


bind()

// Bind socket to port

// We'll fill in this datastruct

struct sockaddr_in socketInfo;

// Fill socketInfo with 0's

memset(&socketInfo, '\0', sizeof(socketInfo));

// Use std TCP/IP

socketInfo.sin_family = AF_INET;

// Tell port in network endian with htons()

socketInfo.sin_port = htons(portNumber);

// (1) Allow connections from same machine only:

struct in_addr addr;

if (inet_aton("127.0.0.1", &addr) == 0) exit(EXIT_FAILURE);

socketInfo.sin_addr.s_addr = addr.s_addr;

// or (2) Allow machine to connect to this service

socketInfo.sin_addr.s_addr = INADDR_ANY;

// Try to bind socket with port and other specifications

int status = bind(socketDescriptor, // from socket()

(struct sockaddr*)&socketInfo,

sizeof(socketInfo));

status == -1 on error

bind()

```
// Bind socket to port
```

```
// We'll fill in this datastruct
```

```
struct sockaddr_in socketInfo;
```

```
// Fill socketInfo with 0's
```

```
memset(&socketInfo, '\0', sizeof(socketInfo));
```

```
// Use std TCP/IP
```

```
socketInfo.sin_family = AF_INET;
```

```
// Tell port in network endian with htons()
```

```
socketInfo.sin_port = htons(portNumber);
```

```
// (1) Allow connections from same machine only:
```

```
struct in_addr addr;
```

```
if (inet_aton("127.0.0.1", &addr) == 0) exit(EXIT_FAILURE);
```

```
socketInfo.sin_addr.s_addr = addr.s_addr;
```

```
// Try to bind socket with port and other specifications
```

```
int status = bind(socketDescriptor, // from socket()
```

```
                (struct sockaddr*)&socketInfo,
```

```
                sizeof(socketInfo));
```

```
status == -1 on error
```

bind()

// Bind socket to port

// We'll fill in this datastruct

struct sockaddr_in socketInfo;

// Fill socketInfo with 0's

memset(&socketInfo, '\0', sizeof(socketInfo));

// Use std TCP/IP

socketInfo.sin_family = AF_INET;

// Tell port in network endian with htons()

socketInfo.sin_port = htons(portNumber);

// or (2) Allow machine to connect to this service

socketInfo.sin_addr.s_addr = INADDR_ANY;

// Try to bind socket with port and other specifications

int status = bind(socketDescriptor, // from socket()

(struct sockaddr*)&socketInfo,

sizeof(socketInfo));

status == -1 on error

bind()

// Bind socket to port

// We'll fill in this datastruct

struct sockaddr_in socketInfo;

// Fill socketInfo with 0's

memset(&socketInfo, '\0', sizeof(socketInfo));

// Use std TCP/IP

socketInfo.sin_family = AF_INET;

// Tell port in network endian with htons()

socketInfo.sin_port = htons(portNumber);

// (1) Allow connections from same machine only:

struct in_addr addr;

if (inet_aton("127.0.0.1", &addr) == 0) exit(EXIT_FAILURE);

socketInfo.sin_addr.s_addr = addr.s_addr;

// or (2) Allow machine to connect to this service

socketInfo.sin_addr.s_addr = INADDR_ANY;

// Try to bind socket with port and other specifications

int status = bind(socketDescriptor, // from socket()

(struct sockaddr*)&socketInfo,

sizeof(socketInfo));

status == -1 on error

What are those structs?

```
typedef uint32_t in_addr_t;
```

```
struct in_addr  
{  
    in_addr_t s_addr;  
};
```

```
struct sockaddr_in  
{  
    sa_family_t    sin_family; // addr family: AF_INET  
    in_port_t      sin_port;   // port (in network  
                                // byte order)  
    struct in_addr sin_addr;    // internet addr  
};
```

listen()

- (Should be called “setQueueSize()”)
- Tell OS how many clients may queue up while server busy


```
int status =  
    listen(socketDescriptor,  
           maxNumPendingClients);
```

- (Almost) ready to listen to port!
- 8 is a good default for maxNumPendingClients.
- If status== -1 then error

Basic accept()

```
// Accept connection to client
int clientDescriptor =
    accept(socketDescriptor, NULL, NULL);
```

- Wait (by default) for someone to actually connect
- Returns
 - a file descriptor for talking with one particular client, or
 - -1 for error
- `connectionDescriptor` for talking with that one client (there may be others for other clients)
- `socketDescriptor` is for listening to socket.



**Waiting
for a client
to call**



Intermediate accept()

Say! Just who are these clients?

```
struct sockaddr_in  clientAddr;  
socklen_t  clientAddrLen = sizeof(clientAddr); //Init to clientAddr  
len  
int clientFd = accept(listenFd,  
    (sockaddr*)&clientAddr, &clientAddrLen);  
  
if (clientFd < 0)  
{  
    perror("accept()");  
    exit(EXIT_FAILURE);  
}  
  
printf("Server: connect from host %s, port %d.\n",  
    inet_ntoa(clientAddr.sin_addr),  
    (int)ntohs(clientAddr.sin_port)  
);
```


Advanced accept()

- Do accept() only when necessary!
- Bust out of accept() after specified time!
 - Can have accept()-ing thread check other things
 - Like flag for “*should I still be running?*”
- **int select(int nfd, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);**
 - **readfds**: set of file descriptors to watch for input
 - **writefds**: set of file descriptors to watch to see if output will not block
 - **exceptfds**: set of file descriptors to watch for exceptions
 - **nfd**: highest number file descriptor in all 3 sets, plus 1.
 - **timeout**: timeout time (returns -1 and sets errno == ?)
- On exit, sets are modified to “these are the file descriptors to handle”
 - Therefore, have to set it back before next select()
- If set is NULL then, “Do not care about these file descriptors”

Advanced accept()

- `void FD_ZERO(fd_set *set);`
 - Clears (initializes) set of file descriptors
- `void FD_SET(int fd, fd_set *set);`
 - Sets 'fd' on in set
- `void FD_CLR(int fd, fd_set *set);`
 - Turn 'fd' off in set
- `int FD_ISSET(int fd, fd_set *set);`
 - Tells if 'fd' is on in given set
- `struct timeval`
 - {
 - long tv_sec; /* seconds */
 - long tv_usec; /* microseconds */

Advanced accept() Example usage

```
struct fd_set originalSet;
struct fd_set modifiedSet;
struct timeval time; time.tv_sec = 1; time.tv_usec = 0;
int listenFd = getListenFd(port);

FD_ZERO(&originalSet);
FD_SET(listenFd,&originalSet);
while (shouldRun)
{
    modifiedSet = originalSet;
    if (select(listenFd+1,&modifiedSet,NULL,NULL,&time) < 0)
    {
        if (errno == EINTR)
            // No biggie, just SIGCHLD or SIGINT or something
        else
            // Ruh-roh! More serious
    }
    if ( !FD_ISSET(listenFd,&modifiedSet) )
        continue; // Probably just timed out
    int clientFd = accept(listenFd, (sockaddr*)&clientAddr,&clientAddrLen);
    // Etc. etc.
}
```

Advanced accept()

```
// From https://www.gnu.org/software/libc/manual/html_node/Server-Example.html
#include <stdio.h>
#include <errno.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#define PORT 5555
#define MAXMSG 512

int read_from_client (int filedес)
{
    char buffer[MAXMSG];
    int nbytes;
    nbytes = read (filedes, buffer, MAXMSG);
    if (nbytes < 0)
    {
        /* Read error. */
        perror ("read");
        exit (EXIT_FAILURE);
    }
    else if (nbytes == 0)
        /* End-of-file. */
        return -1;
    else
    {
        /* Data read. */
        fprintf (stderr, "Server: got message: `%s'\n", buffer);
        return 0;
    }
}
```

```
int main (void)
{
    extern int make_socket (uint16_t port);
    int sock;
    fd_set active_fd_set, read_fd_set;
    int i;
    struct sockaddr_in clientname;
    size_t size;

    /* Create the socket and set it up to accept
connections. */
    sock = make_socket (PORT);
    if (listen (sock, 1) < 0)
    {
        perror ("listen");
        exit (EXIT_FAILURE);
    }

    /* Initialize the set of active sockets. */
    FD_ZERO (&active_fd_set);
    FD_SET (sock, &active_fd_set);
```

Advanced accept()

```
while (1)
{
    /* Block until input arrives on one or more active sockets.
    */
    read_fd_set = active_fd_set;
    if (select (FD_SETSIZE, &read_fd_set, NULL, NULL,
    NULL) < 0)
    {
        perror ("select");
        exit (EXIT_FAILURE);
    }

    /* Service all the sockets with input pending. */
    for (i = 0; i < FD_SETSIZE; ++i)
        if (FD_ISSET (i, &read_fd_set))
        {
            if (i == sock)
            {
                /* Connection request on original socket. */
                int new;
                size = sizeof (clientname);
                new = accept (sock,
                            (struct sockaddr *) &clientname,
                            &size);
```

```
                if (new < 0)
                {
                    perror ("accept");
                    exit (EXIT_FAILURE);
                }
                fprintf
                (stderr,
                 "Server: connect from host %s, port %hd.\n",
                 inet_ntoa (clientname.sin_addr),
                 (int)ntohs (clientname.sin_port));
                FD_SET (new, &active_fd_set);
            }
        }
    else
    {
        // Data arriving on an already-connected socket.
        if (read_from_client (i) < 0)
        {
            close (i);
            FD_CLR (i, &active_fd_set);
        }
    }
}
```

More C socket goodies

- Just like Java, C lets you get and set socket-related parameters
 - `int getsockopt(int sockfd, int level, int optname, void *optval, socklen_t *optlen);`
 - `int setsockopt(int sockfd, int level, int optname, const void *optval, socklen_t optlen);`

And thread pools in C/C++? (1)

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>

const int  NUM_THREADS = 16;

const int  BUFFER_SIZE  = 16;

static
bool shouldRun__ = true;

bool getShouldRun ()
{
    return(shouldRun__);
}

class Buffer
{
    // I. Member vars:
    int  array_[BUFFER_SIZE];

    size_t inIndex_;

    size_t outIndex_;

    size_t count_;
```

```
pthread_mutex_t  lock_;

pthread_cond_t   notEmpty_;

// II. Disallowed auto-generated methods:
// No copy constructor:
Buffer          (const Buffer&);

// No copy assignment op:
Buffer& operator= (const Buffer&);

protected :
    // III. Protected methods:

public :
    // IV. Constructor(s), assignment op(s), factory(s) and destructor:
    Buffer          () :
        inIndex_(0),
        outIndex_(0),
        count_(0)
    {
        pthread_mutex_init(&lock_,NULL);
        pthread_cond_init(&notEmpty_,NULL);
    }

    // PURPOSE To release the resources of '*this'. No parameters.
    //No return value.
    ~Buffer          ()
    {
        pthread_cond_destroy(&notEmpty_);
        pthread_mutex_destroy(&lock_);
    }
```

And thread pools in C/C++? (2)

```
// V. Accessors:

// VI. Mutators:

// VII. Methods that do main and misc. work of class:
// PURPOSE: To insert file descriptor 'fd', address 'addr' and the
// length of the data in 'addr' noted in 'addrLen' in '*this',
// *ONLY* if there is space. No return value.
void put (int value
)
{
pthread_mutex_lock(&lock_);

if (count_ < BUFFER_SIZE)
{
array_[inIndex_] = value;

if (++inIndex_ >= BUFFER_SIZE)
inIndex_ = 0;

count_++;
pthread_cond_signal(&notEmpty_);
}

pthread_mutex_unlock(&lock_);
}
```

```
// PURPOSE: To wait until data is available, and then to set 'value' equal
//to the next available values stored in '*this'. No return value.
void get (int id,
int& value
)
{
pthread_mutex_lock(&lock_);

while ( getShouldRun() &&
(count_ == 0)
)
{
pthread_cond_wait(&notEmpty_,&lock_);

if ( !getShouldRun() )
{
printf("Ending worker thread %d\n",id);
pthread_mutex_unlock(&lock_);
pthread_exit(NULL);
}

}

value = array_[outIndex_];

if (++outIndex_ >= BUFFER_SIZE)
outIndex_ = 0;

count_--;
pthread_mutex_unlock(&lock_);
}
```


And thread pools in C/C++? (3)

```
// PURPOSE: To notify all waiting threads to wake up. No parameters.
// No return value.
void wakeWaiters ()
{
    pthread_mutex_lock(&lock_);
    pthread_cond_broadcast(&notEmpty_);
    pthread_mutex_unlock(&lock_);
}
};
```

```
Buffer buffer;
```

```
void* obtainerThreadFnc (void* vPtr)
{
    int i;
    char line[BUFFER_SIZE];

    while (getShouldRun())
    {
        printf("Enter an integer, or to 0 quit: ");
        fflush(stdout);
        fgets(line,BUFFER_SIZE,stdin);
        i = strtol(line,NULL,0);

        if (i == 0)
            break;

        buffer.put(i);
    }

    shouldRun__ = false;
    printf("Ending obtaining thread.\n");
}
```

```
void* workerThreadFnc (void* vPtr
)
{
    int id = *(int*)vPtr;
    int value;

    printf("Thread %d \"Ready for work!\"\n",id);
    fflush(stdout);

    while (getShouldRun())
    {
        buffer.get(id,value);

        printf("Thread %d: \"%2 * %d = %d\"\n",id,value,2*value);
        fflush(stdout);
    }

    return(NULL);
}

void initializeThreads (pthread_t* obtainerThreadPtr,
pthread_t* workerThreadArray
)
{
    int idInfo[NUM_THREADS];

    for (int i = 0; i < NUM_THREADS; i++)
    {
        idInfo[i] = i;
        pthread_create(workerThreadArray+i,NULL,workerThreadFnc,(void*)(idInfo+i));
    }

    pthread_create(obtainerThreadPtr,NULL,obtainerThreadFnc,NULL);
}
```

And thread pools in C/C++? (3)

```
int main ()
{
    pthread_t  obtainerThread;
    pthread_t  workerThreadArray[NUM_THREADS];

    initializeThreads(&obtainerThread,workerThreadArray);

    while ( getShouldRun() )
        sleep(1);

    shouldRun__ = false;
    buffer.wakeWaiters();

    for (int index = 0; index < NUM_THREADS; index++)
        pthread_join(workerThreadArray[index],NULL);

    pthread_join(obtainerThread,NULL);

    return(EXIT_SUCCESS);
}
```

Your turn!

C would benefit from thread pools too.

Write a multi-threaded uppercasing server in C/C++ that:

- (1) Uses a thread pool
- (2) Records clients IPs
- (3) Times-out after 1 sec/sets a global flag to tell when still running

References:

- Elliotte Rusty Harold “*Java Network Programming: 4th Ed.*”
- M. Tim Jones “*Gnu/Linux Application Programming, 2nd Ed.*” Charles River Media/Course Technology Cengage Learning.
- https://www.gnu.org/software/libc/manual/html_node/Server-Example.html