

# Distributed Systems

## Lecture 10

### Distributing a System

Joseph Phillips

Copyright (c) 2019

Last modified 2019 December 29

# Reading


- “*Designing Distributed Systems: Patterns and Paradigms for Scalable, Reliable Services*” by Brendan Burns, O’Reilly

# Topics

- Distributing a System

We have a system!  
It is very popular!

Clients



Gimme!  
Gimme!  
Gimme!

Server




Whoa! Slowdown!  
One-at-a-time!

Time to scale up!

*If only we knew how?*

Hey! Let us consider design-patterns for distributed systems!

# Sidecar Design Pattern

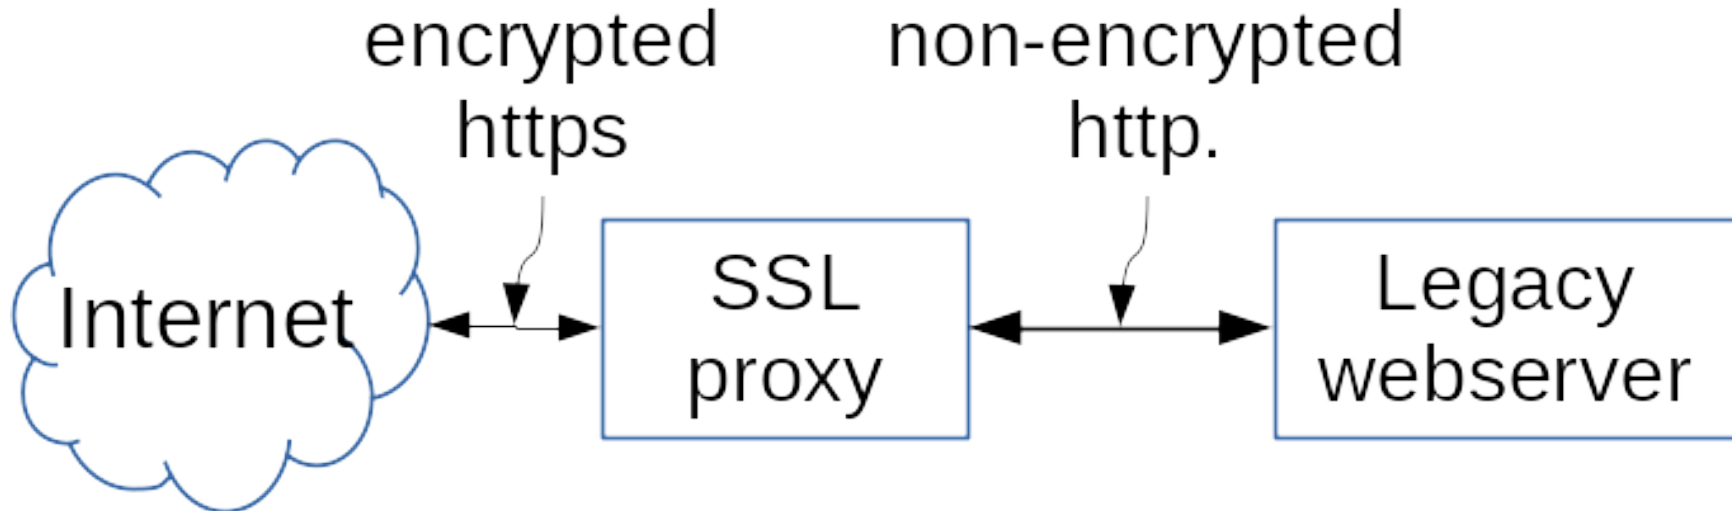


- Goal: Improve functionality of application (without app knowing it)




# Side-car in use

- Example: Add SSL front end to legacy server



# Side-car in use

- Example: Allow dynamic configuration of legacy application




- 1) Configuration process continually checks cloud-based config files
- 2) When change noted, it writes to legacy file in legacy format
- 3) Config process then notifies (e.g. signals or restarts) application
- 4) App reloads legacy file



# Ambassador Design Pattern

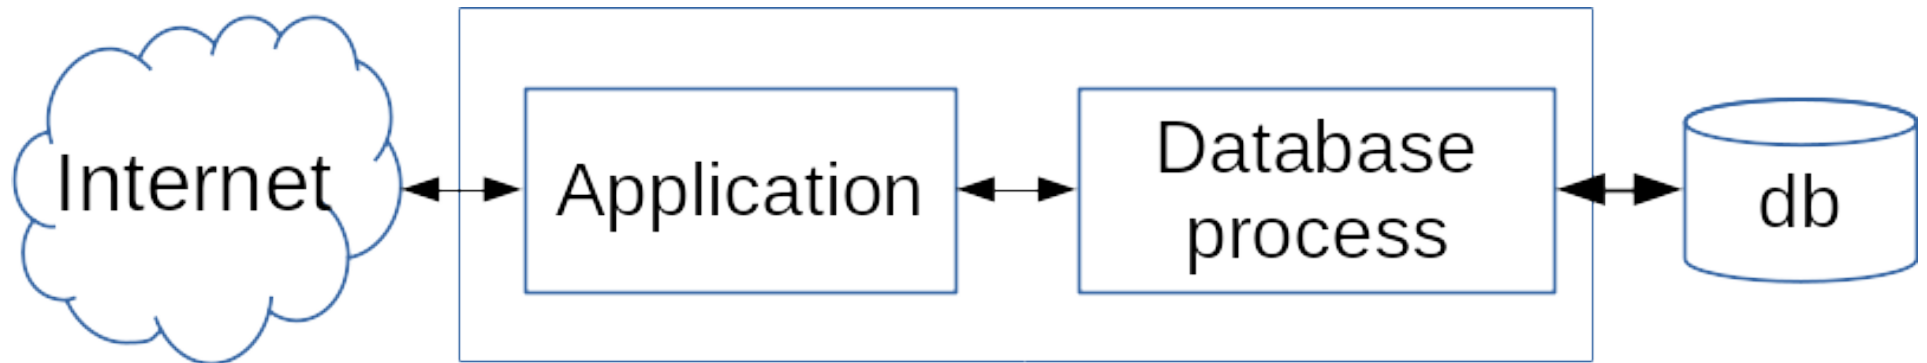


- Goal: To broker between application and other resources it might need




# Ambassador in use

- Example: Database
  - Prototypical example of ambassador pattern




# Ambassador in use

- Example: Multi-sharded application
  - Different shards have different responsibilities
  - E.g. handle different users




# Serving Patterns

- We have 2 basic design patterns:
  - sidecar (app behind some helper process)
  - ambassador (app in front of some helper process)
- Let us apply them to larger, more realistic, systems
- ***But first***, how shall we even design a larger system?




- The *Monolithic* design
  - Any problems?

# Serving Patterns



- The *Microservice* design
- More reliable
  - If one thing crashes, does not bring down everything
- More agile
  - Need more of one particular service? **No problem!**
  - Make more processes of just that desired microservice

# Load balanced Architecture



- Load balancer is an application of ***sidecar***
- Which replica gets which request?
  - Need a hash-function
- What about IP-address?
  - Works okay for ***intra***net apps
  - Fails for ***Inter***net apps: IP addresses too clumped
- Consistent hash function
  - Allows adding/removing replicas as become more/less busy

# Consistent hash function


- Motivation:
  - We dynamically add and remove servers
  - We do not want to drastically alter the hash function of clients to servers when we do so
  - (Thanks to Juan Pablo Carzollo, who attributes Karger et al from MIT)

# InConsistent hash function

- $\text{serverToUse} = \text{userHash}(\text{user}) \% \text{numServers}$
- 3 servers:
  - huey (0), dewey (1), louie (2)
- 4 users:
  - Alice (hash=50), Bob (hash=150), Cathy (hash=250), David (hash=350)
- Original mapping:
  - Alice:  $50 \% 3 = 2$  (louie)
  - Bob:  $150 \% 3 = 0$  (huey)
  - Cathy:  $250 \% 3 = 1$  (dewey)
  - David:  $350 \% 3 = 2$  (louie)



# InConsistent hash function


- Now remove a server:
  - $\text{serverToUse} = \text{userHash}(\text{user}) \% \text{numServers}$
- 2 servers:
  - huey (0), dewey (1)
- Many changes to the hash fnc! 

- Original mapping:
  - Alice:  $50 \% 3 = 2$  (louie)
  - Bob:  $150 \% 3 = 0$  (huey)
  - Cathy:  $250 \% 3 = 1$  (dewey)
  - David:  $350 \% 3 = 2$  (louie)

- New mapping:
  - Alice:  $50 \% 2 = 0$  (**huey**)
  - Bob:  $150 \% 2 = 0$  (huey)
  - Cathy:  $250 \% 2 = 0$  (**huey**)
  - David:  $350 \% 2 = 0$  (**huey**)


# Consistent Hash Function (0)

- Consider a circle from 0 to 360 degrees
- Give each server several (in this case 4) random positions on the circle
- The 4 ***is this example only!*** In real life, you would have more!




# Consistent Hash Function (1)

- Now hash the users on the same circle:
  - $\text{userHash}(\text{user}) \% 360$




# Consistent Hash Function (2)

- Map each user to a server by (for example) going counter-clockwise
  - Alice => Louie
  - Bob => Dewey
  - Cathy => Dewey
  - David => Louie




# Consistent Hash Function (3)

- Remove a server!
- Only approx.  $\frac{1}{2}$  remaining mappings change:
  - Alice => **Huey**
  - Bob => Dewey
  - Cathy => Dewey
  - David => **Dewey**



# Consistent Hash Function (4)

- Add a server!
- Only approx. 1/3 existing mappings change:
  - Alice => **Daisy**
  - Bob => Dewey
  - Cathy => Dewey
  - David => Dewey



# Consistent Hash Function (5)

Q: “Hey, server Daisy is *twice as powerful* as Huey, Dewey or Louie!”

A: “Then give Daisy *twice as many* random points on the circle”

# Consistent Hash Function (6)

- We will randomly spread points over a circle
- But what should ***userHash()*** be base on?
- Well, an HTTP request has several things including:
  - time (e.g. 12:00:01)
  - IP address (e.g. 1.2.3.4)
  - path (e.g. /directory/page.html)



# Consistent Hash Function (7)


- So, let us consider them:
- Time
  - *Too specific!* 12:00:01 different from 12:00:02
- IP address
  - *Too specific!* 1.2.3.4 different from everything
- Path
  - *Good idea!* Go right to /directory/page.html

# Consistent Hash Function (8)



- One teeny, tiny technicality . . . display these differently:
  - /directory/page.html (LANG=en\_US)
  - /directory/page.html (LANG=es\_MX)
- What Brendan Burns recommends:
  - ***hash(country(request.ipAddr), request.path)***

# Load balanced Architecture

- Variations on a theme
  - Load balanced and cached architecture
- Look ma! ***Double sidecar!***
  - Cache's relationship with balancer
  - Replica's relationship with cache



# Replicas vs. Shards



- Replicas:

- All requests routed to all replicas
- Gives **redundancy**

- Shards:

- Specific requests routed to specific shards
- Gives **efficiency**


You can, of course, **do both**  
at different scales and levels!

# Scatter/Gather Design Pattern

- Motivation:
  - Split requests up into multiple sub-requests
  - Give each sub-request to its own shard
- ***Look Ma!*** This is the ***ambassador design pattern*** applied multiple times


# Scatter/Gather (0)

- Client “*Hey server: give me info about **chocolate** and **vanilla**!*”




# Scatter/Gather (1)

- Router/Combiner “*Shard0, you handle **chocolate**. Shard1, you handle **vanilla**.*”



# Scatter/Gather (2)


- Shard0 and Shard1 “*Here you go!*”





# Scatter/Gather (3)

- Router/Combiner “*Chocolate and Vanilla, as you requested*”



# References:

- Brendan Burns “*Designing Distributed Systems: Patterns and Paradigms for Scalable, Reliable Services*”
- <https://www.geeksforgeeks.org/parse-json-java/>
- [https://www.w3schools.com/js/js\\_json\\_xml.asp](https://www.w3schools.com/js/js_json_xml.asp)
- <https://www.json.org/>
- “The Ultimate Guide to Hashing Content”  
<https://www.toptal.com/big-data/consistent-hashing> Juan Pablo Carzollo, downloaded 2019 Nov 20
- Karger, David; Lehman, Eric; Leighton, Tom; Levine, Matthew; Lewin, Daniel; Panigrahy, Rina “Consistent Hashing and Random Trees: Distributed Caching Protocols for Relieving Hot Spots on the World Wide Web”